

A word is worth a thousand vectors

(word2vec, lda, and introducing **lda2vec**)

Christopher Moody
@ Stitch Fix



Welcome,
thanks for coming, having me, organizer

NLP can be a messy affair because you have to teach a computer about the irregularities and ambiguities of the English language in this sort of hierarchical sparse nature in all the grammar

3rd trimester, pregnant

“wears scrubs” — medicine

taking a trip — a fix for vacation clothing

power of word vectors promise is to sweep away a lot of issues

About




 [@chrisemoody](https://twitter.com/chrisemoody)

 Caltech Physics

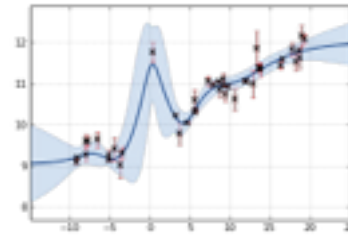
 PhD. in astrostats supercomputing

 sklearn t-SNE contributor

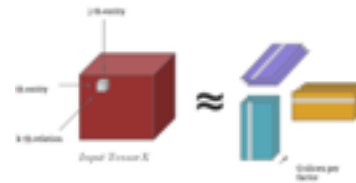
 Data Labs at Stitch Fix

github.com/cemoody

Gaussian Processes



Tensor Decomposition



t-SNE



chainer
deep learning



Credit

Large swathes of this talk are from
previous presentations by:

- [Tomas Mikolov](#)
- [David Blei](#)
- [Christopher Olah](#)
- [Radim Rehurek](#)
- [Omer Levy & Yoav Goldberg](#)
- [Richard Socher](#)
- [Xin Rong](#)
- [Tim Hopper](#)

1 word2vec

2 lda

3 lda2vec

word2vec

1. *king - man + woman = queen*
2. Huge splash in NLP world
3. Learns from raw text
4. Pretty simple algorithm
5. Comes pretrained

1. Learns what words *mean* — can solve analogies cleanly.
 1. Not treating words as blocks, but instead modeling relationships
2. Distributed representations form the basis of more complicated deep learning systems
3. Shallow — not deep learning!
 1. Power comes from this simplicity — super fast, lots of data
4. Get a lot of mileage out of this
 1. Don't need to model the wikipedia corpus before starting your own

word2vec

1. Set up an objective function
2. Randomly initialize vectors
3. Do gradient descent

word2vec

word2vec: learn word vector v_{in}
from it's surrounding context

v_{in}

1. Let's talk about training first
2. In SVD and n-grams we built a co-occurrence and transition probability matrices
3. Here we will learn the embedded representation directly, with no intermediates, update it w/ every example

“The fox jumped **over** the lazy dog”

Maximize the likelihood of seeing the words given the word **over**.

$$\begin{aligned}P(the|over) \\ P(fox|over) \\ P(jumped|over) \\ P(the|over) \\ P(lazy|over) \\ P(dog|over)\end{aligned}$$

...instead of maximizing the likelihood of co-occurrence counts.

1. Context — the words surrounding the training word
2. Naively assume $P(*|over)$ is independent conditional on the training word
3. Still a pretty simple assumption!

Conditioning on just *over* no other secret parameters or anything

word2vec

What should this be?

$$P(fox|over)$$

word2vec

Should depend on the word vectors.

$$P(fox|over)$$



$$P(v_{fox}|v_{over})$$

Trying to learn the word vectors, so let's start with those
(we'll randomly initialize them to begin with)

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped **over** the lazy dog”

word2vec

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped **over** the lazy dog”
↑
 v_{IN}

IN = training word

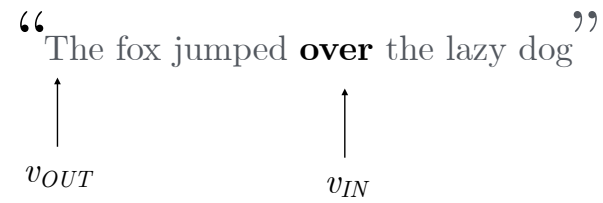
Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped **over** the lazy dog”

v_{OUT} v_{IN}



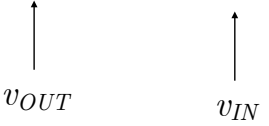
Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped **over** the lazy dog”

v_{OUT} v_{IN}



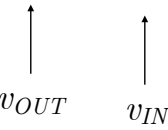
Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped **over** the lazy dog”

v_{OUT} v_{IN}



Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped **over** the lazy dog”

↑ ↑
 v_{IN} v_{OUT}

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped **over** the lazy dog”

v_{IN} v_{OUT}

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped **over** the lazy dog”

v_{IN} v_{OUT}

A red triangle with the text "word2vec" written inside it.

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

\uparrow
 v_{IN}

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”
 ↑ ↑
 v_{OUT} v_{IN}

...So that at a high level is what we want word2vec to do.

word2vec

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

word2vec

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

word2vec

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

word2vec

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

word2vec

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

...So that at a high level is what we want word2vec to do.

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

...So that at a high level is what we want word2vec to do.

word2vec

Twist: we have *two* vectors for every word.

Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

word2vec

Twist: we have *two* vectors for every word.

Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

word2vec

Twist: we have *two* vectors for every word.

Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

word2vec

Twist: we have *two* vectors for every word.

Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

word2vec

Twist: we have *two* vectors for every word.

Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{OUT} v_{IN}

...So that at a high level is what we want word2vec to do.

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”
 ↑ ↑
 v_{IN} v_{OUT}

...So that at a high level is what we want word2vec to do.

word2vec

Twist: we have *two* vectors for every word.
Should depend on whether it's the input or the output.

Also a *context* window around every input word.

$$P(v_{OUT}|v_{IN})$$

“The fox jumped over **the** lazy dog”

v_{IN} v_{OUT}

...So that at a high level is what we want word2vec to do.

two for loops

That's it! A bit disengious to make this a giant network

objective

How should we define $P(v_{OUT}|v_{IN})$?

Measure loss between
 v_{IN} and v_{OUT} ?

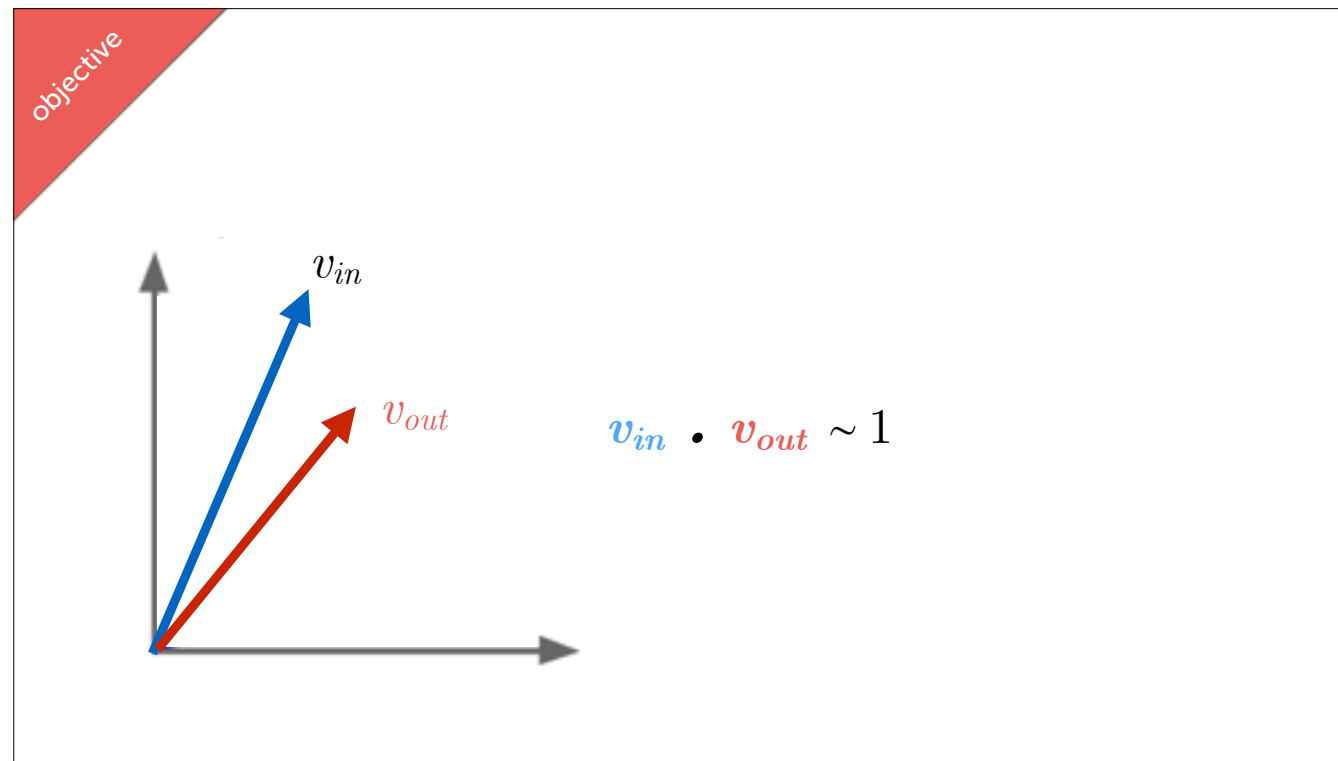
$$v_{in} \cdot v_{out}$$

Now we've defined the high-level update path for the algorithm.

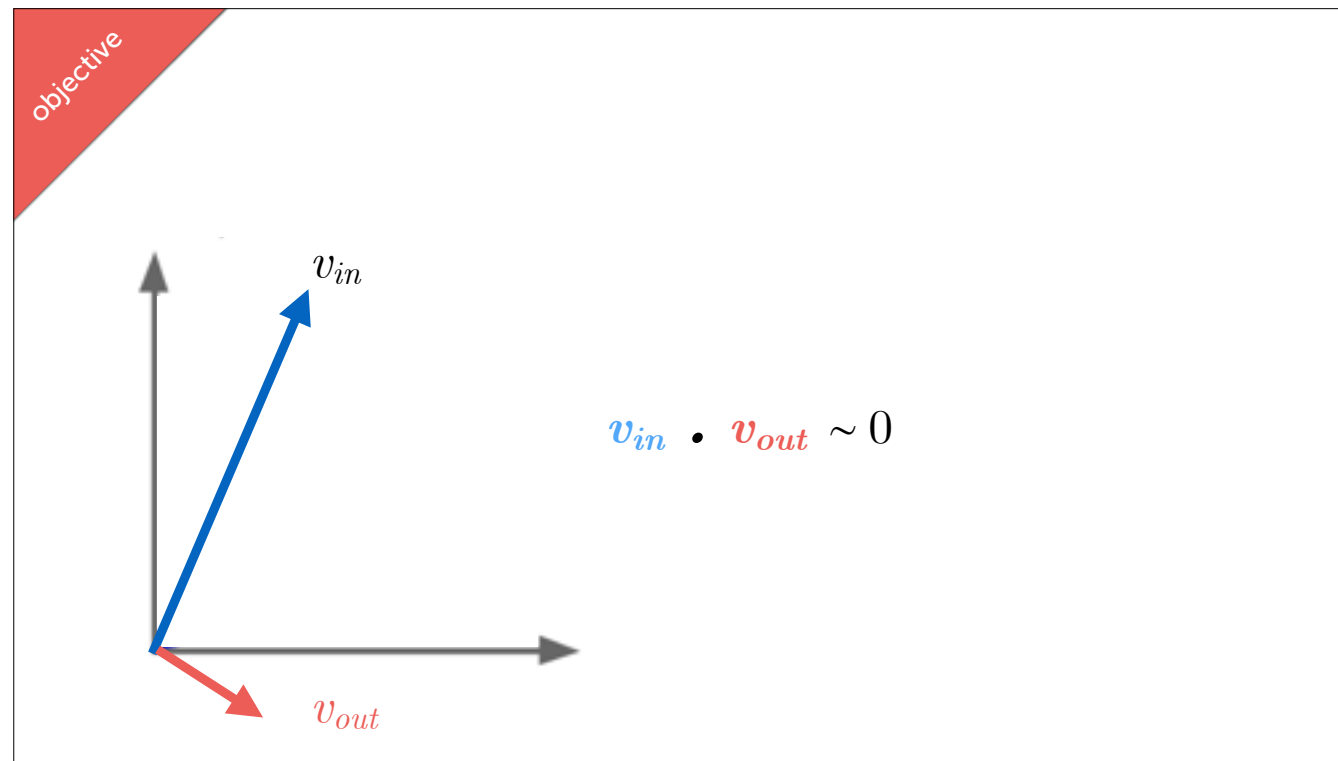
Need to define this prob exactly in order to define our updates.

Boils down to diff between in & out — want to make as similar as possible, and then the probability will go up.

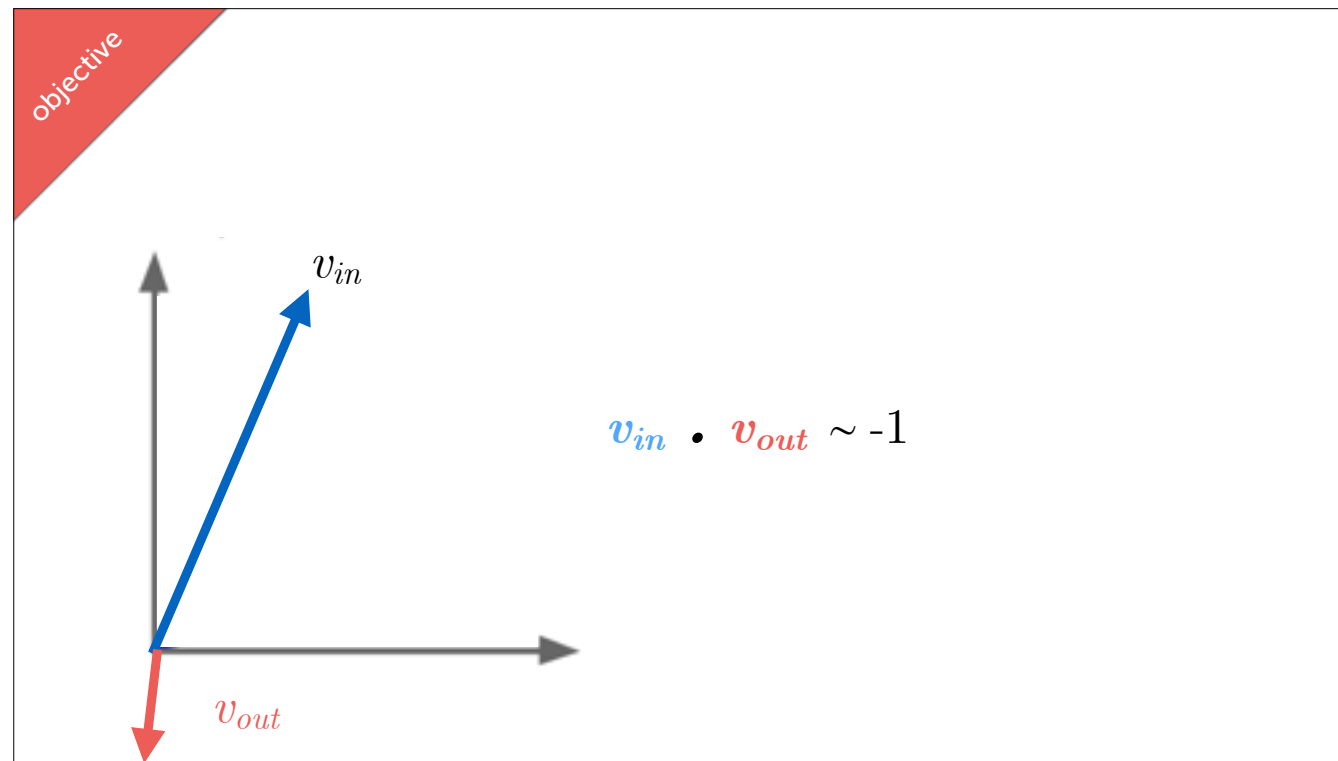
Use cosine sim.



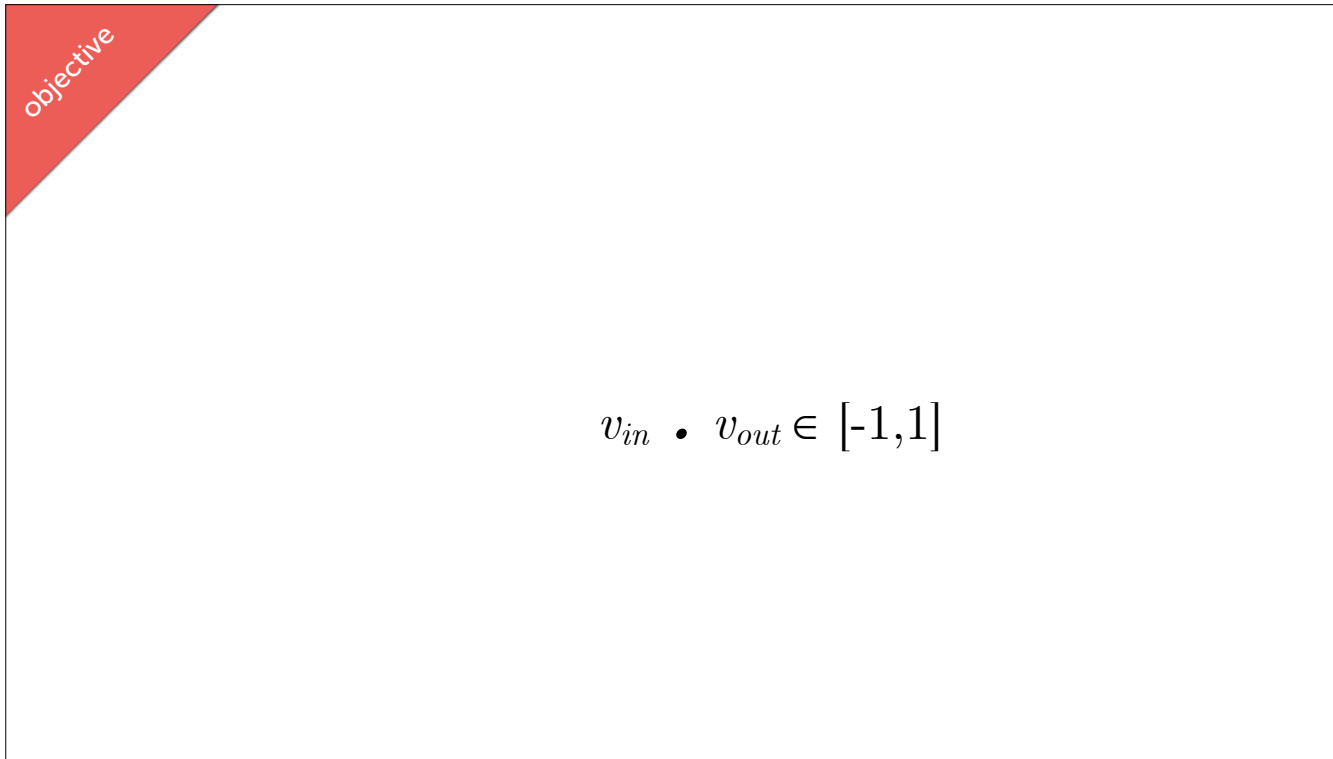
Dot product has these properties:
Similar vectors have similarity near 1



Orthogonal vectors have similarity near 0



Orthogonal vectors have similarity near -1



But the inner product ranges from -1 to 1 (when normalized)
...and we'd like a probability

objective

But we'd like to measure a probability.

$$v_{in} \cdot v_{out} \in [-1, 1]$$

But the inner product ranges from -1 to 1 (when normalized)
...and we'd like a probability

objective

But we'd like to measure a probability.

$$\text{softmax}(v_{in} \cdot v_{out}) \in [0,1]$$

Transform again using softmax

objective

But we'd like to measure a probability.

$$\text{softmax}(v_{in} \cdot v_{out})$$

Probability of choosing 1 of N discrete items.
Mapping from vector space to a multinomial over words.

Similar to logistic function for binary outcomes, but instead for 1 of N outcomes.

So now we're modeling the probability of a word showing up as the combination of the training word vector and the target word vector and transforming it to a 1 of N

But we'd like to measure a probability.

$$\text{softmax} \sim \exp(v_{in} \cdot v_{out}) \in [0,1]$$

So here's the actual form of the equation — we normalize by the sum of all of the other possible pairs of word combinations

objective

But we'd like to measure a probability.

$$\textit{softmax} = \frac{\exp(v_{in} \cdot v_{out})}{\sum_{k \in V} \exp(v_{in} \cdot v_k)}$$

Normalization term over all words

So here's the actual form of the equation — we normalize by the sum of all of the other possible pairs of word combinations

two effects

make v_{in} and v_{out} more similar

make v_{in} and every other word less similar

But we'd like to measure a probability.

$$\text{softmax} = \frac{\exp(v_{in} \cdot v_{out})}{\sum_{k \in V} \exp(v_{in} \cdot v_k)} = P(v_{out} | v_{in})$$

This is the kernel of the word2vec. We're just going to apply this operation every time we want to update the vectors.

For every word, we're going to have a context window, and then for every pair of words in that window and the input word, we'll measure this probability.

objective

Learn by gradient descent on the softmax prob.

For every example we see update v_{in}

$$v_{in} := v_{in} + \frac{\partial}{\partial v_{in}} P(v_{out} | v_{in})$$

$$v_{out} := v_{out} + \frac{\partial}{\partial v_{out}} P(v_{out} | v_{in})$$

...I won't go through the derivation of the gradient, but this is the general idea

relatively simple, fast — fast enough to read billions of words in a day

Model (training time)	Redmond	Havel	ninjutsu
Collobert (50d) (2 months)	conyers lubbock keene	plawen dzerzhinsky osterreich	reiki kohona karate
Turian (200d) (few weeks)	McCarthy Alston Cousins	Jewell Arzu Ovitz	- - -
Mnih (100d) (7 days)	Podhurst Harlang Agarwal	Pontiff Pinochet Rodionov	- - -
Skip-Phrase (1000d, 1 day)	Redmond Wash. Redmond Washington Microsoft	Vaclav Havel president Vaclav Havel Velvet Revolution	ninja martial arts swordsmanship

← word2vec

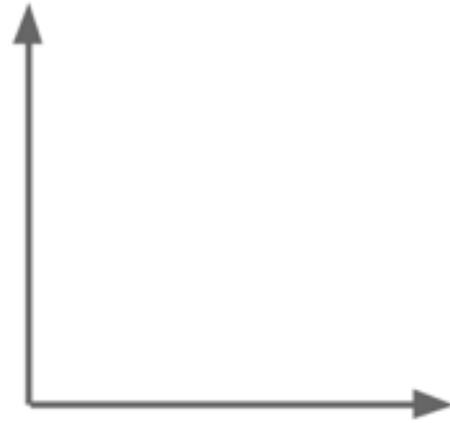
explain table

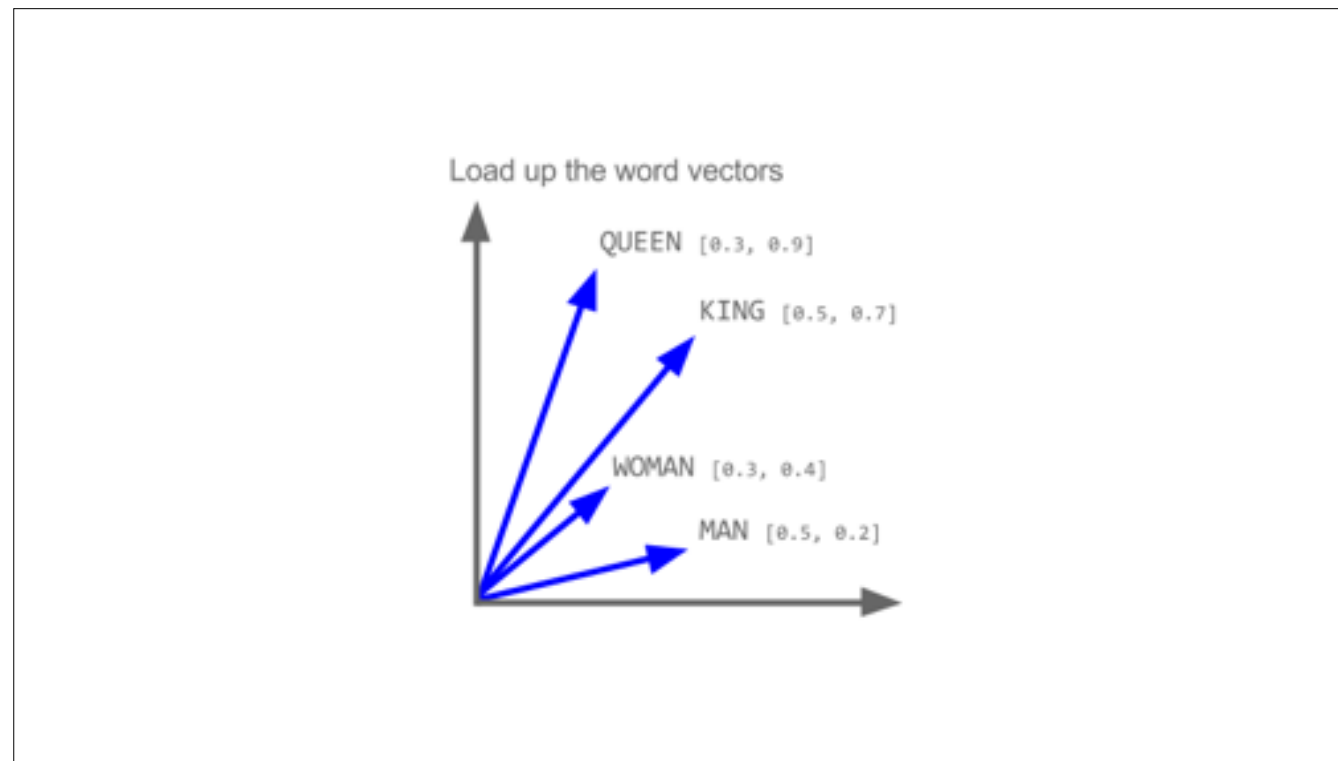
Model	Vector Dimensionality	Training words	Accuracy [%]		
			Semantic	Syntactic	Total
Collobert-Weston NNLM	50	660M	9.3	12.3	11.0
Turian NNLM	50	37M	1.4	2.6	2.1
Turian NNLM	200	37M	1.4	2.2	1.8
Mnih NNLM	50	37M	1.8	9.1	5.8
Mnih NNLM	100	37M	3.3	13.2	8.8
Mikolov RNNLM	80	320M	4.9	18.4	12.7
Mikolov RNNLM	640	320M	8.6	36.5	24.6
Huang NNLM	50	990M	13.3	11.6	12.3
CBOW	300	783M	15.5	53.1	36.1
Skip-gram	300	783M	50.0	55.9	53.3

← word2vec

if not convinced by qualitative results....

What is king + man - woman?





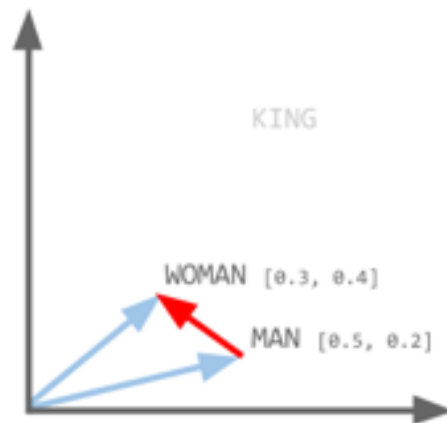
Showing just 2 of the ~500 dimensions. Effectively we've PCA'd it

Start with $\text{man} - \text{woman}$

KING

WOMAN [0.3, 0.4]

MAN [0.5, 0.2]



Start with man - woman

KING

MAN - WOMAN



Then take king

KING [0.5, 0.7]

MAN - WOMAN



A 2D coordinate system is shown with a horizontal x-axis and a vertical y-axis, both ending in arrows. A blue vector originates from the origin (0,0) and points into the first quadrant. A red vector also originates from the origin and points into the first quadrant, but at a shallower angle than the blue vector. The blue vector is labeled 'KING [0.5, 0.7]' and the red vector is labeled 'MAN - WOMAN'.

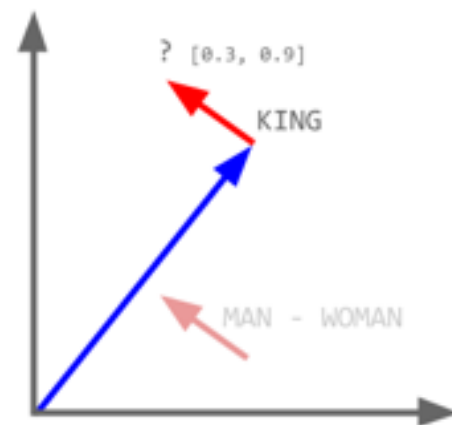
And add $\text{man} - \text{woman}$



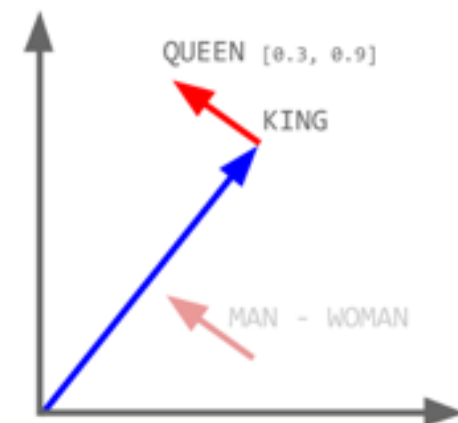
And add $\text{man} - \text{woman}$



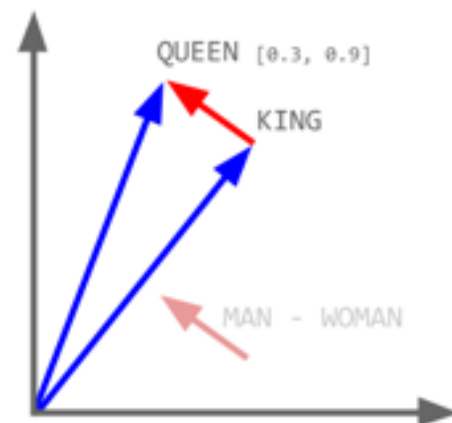
Find nearest word to result



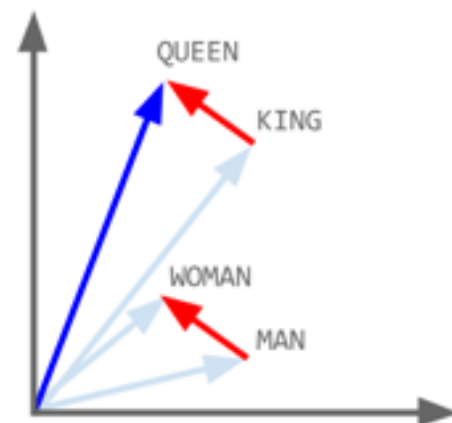
queen is closest to resulting vector



queen is closest to resulting vector

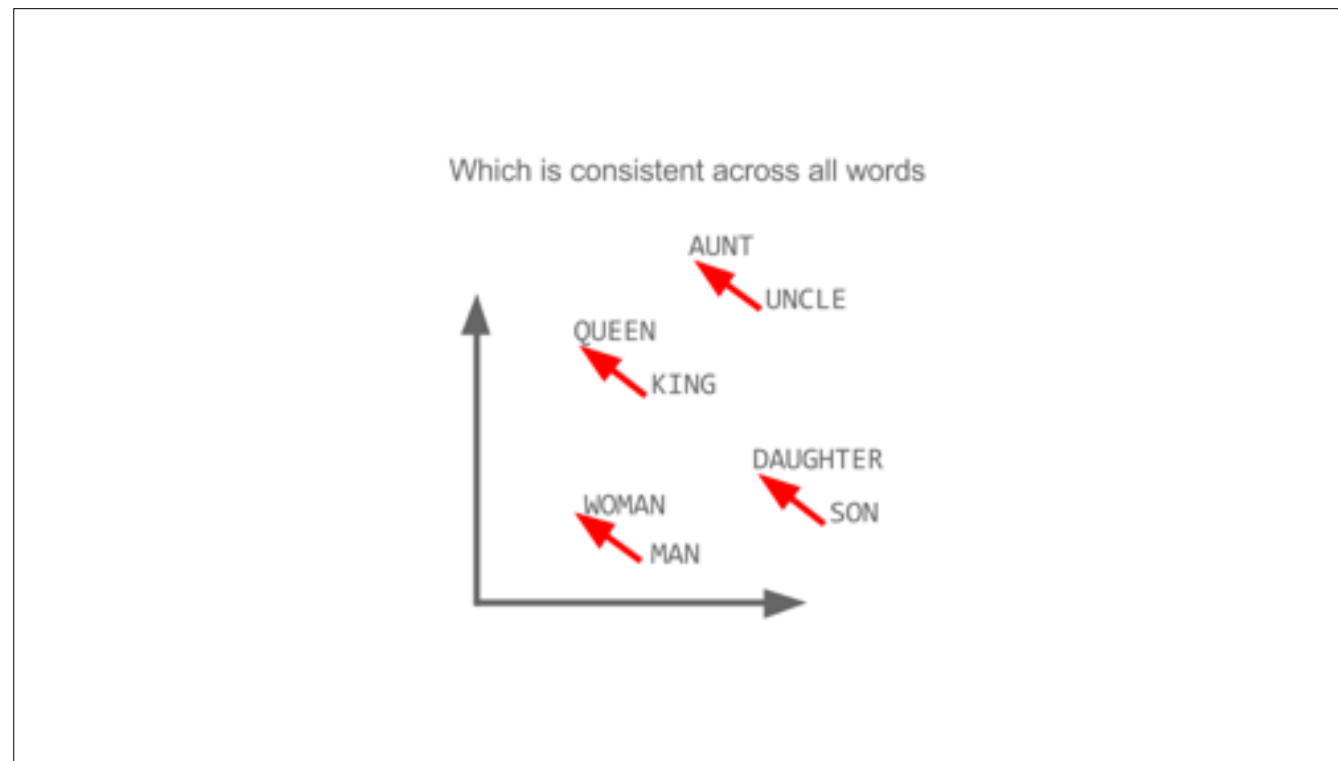


So $\text{king} + \text{man} - \text{woman} = \text{queen!}$



The **red direction** encodes gender





If we only had locality and not regularity, this wouldn't necessarily be true

This **direction** always means **gender**



We have hundreds of **directions**
encoding hundreds of ideas

HIGHER STATUS

MORE FEMININE

Czech + currency	Vietnam + capital	German + airlines	Russian + river	French + actress
koruna	Hanoi	airline Lufthansa	Moscow	Juliette Binoche
Check crown	Ho Chi Minh City	carrier Lufthansa	Volga River	Vanessa Paradis
Polish zolty	Viet Nam	flag carrier Lufthansa	upriver	Charlotte Gainsbourg
CTK	Vietnamese	Lufthansa	Russia	Cecile De

So we live in a vector space where operations like addition and subtraction are meaningful.

So here's a few examples of this working.

Really get the idea of these vectors as being 'mixes' of other ideas & vectors

ITEM_3469 + 'Pregnant'

SF is a person service

Box



+ 'Pregnant'

I love the stripes and the cut around my neckline was amazing

someone else might write 'grey and black'

subtlety and nuance in that language

We have lots of this interaction — of order wikipedia amount — far too much to manually annotate anything

= ITEM_701333
= ITEM_901004
= ITEM_800456

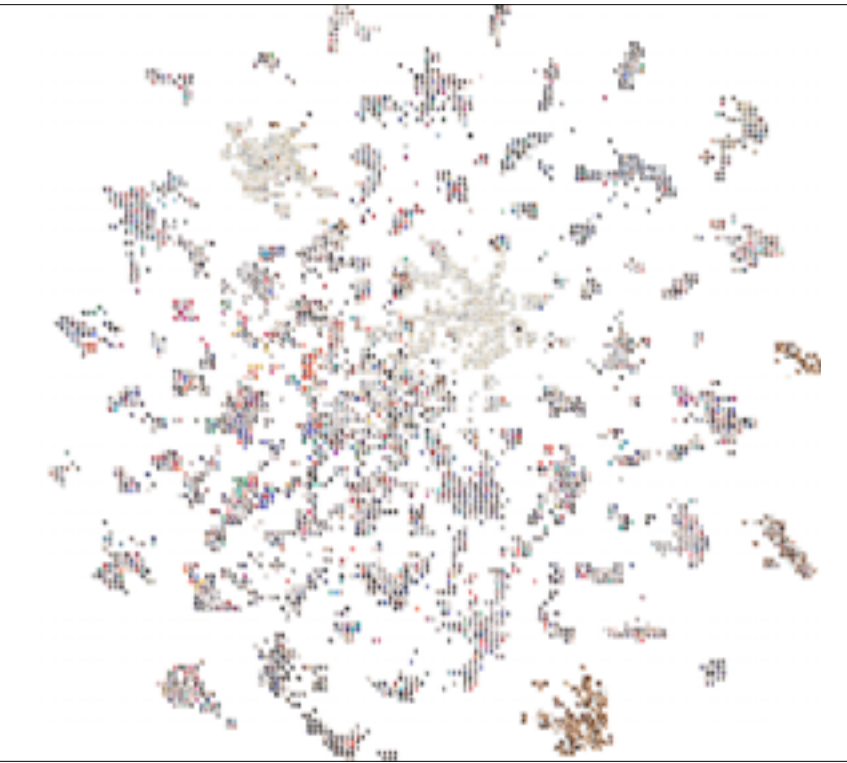


Stripes and are safe for maternity

And also similar tones and flowy — still great for expecting mothers

what about **LDA**?

LDA
on Client Item
Descriptions

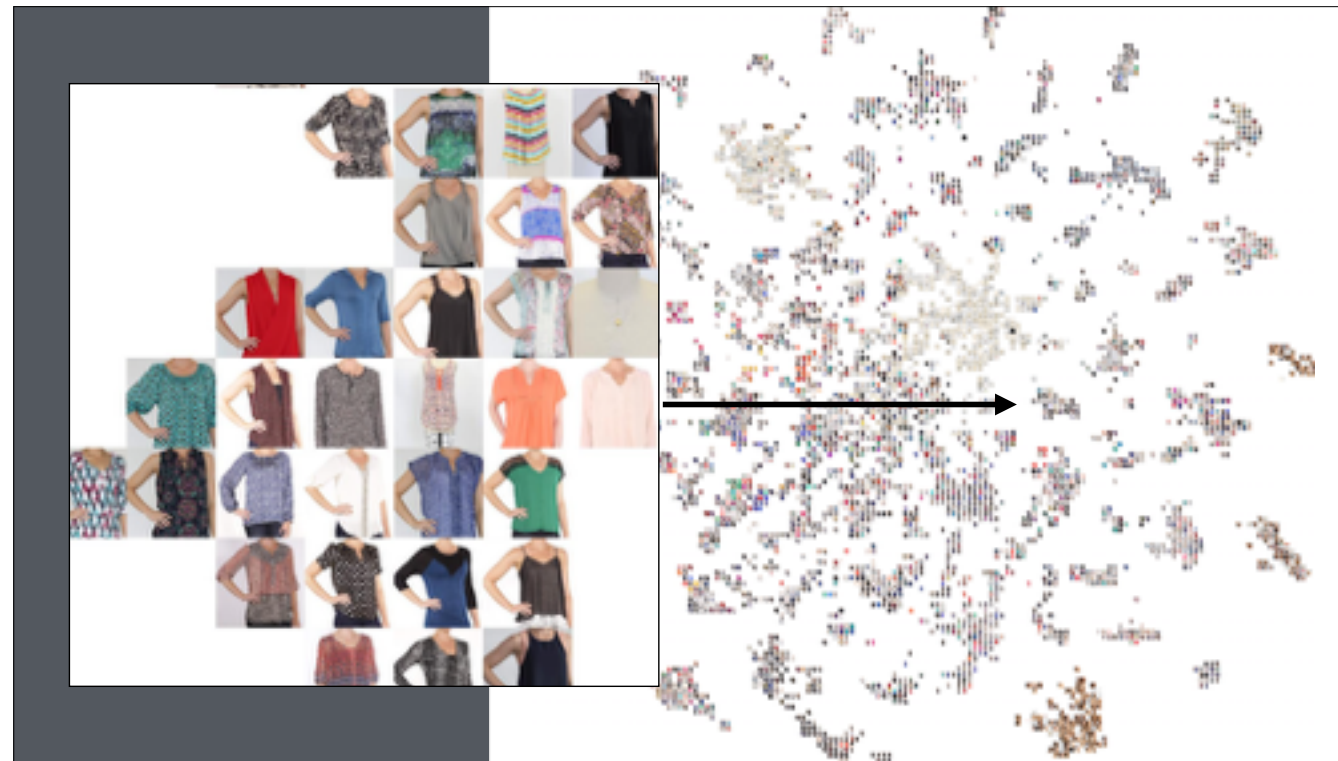


This shows the incredible amount of structure

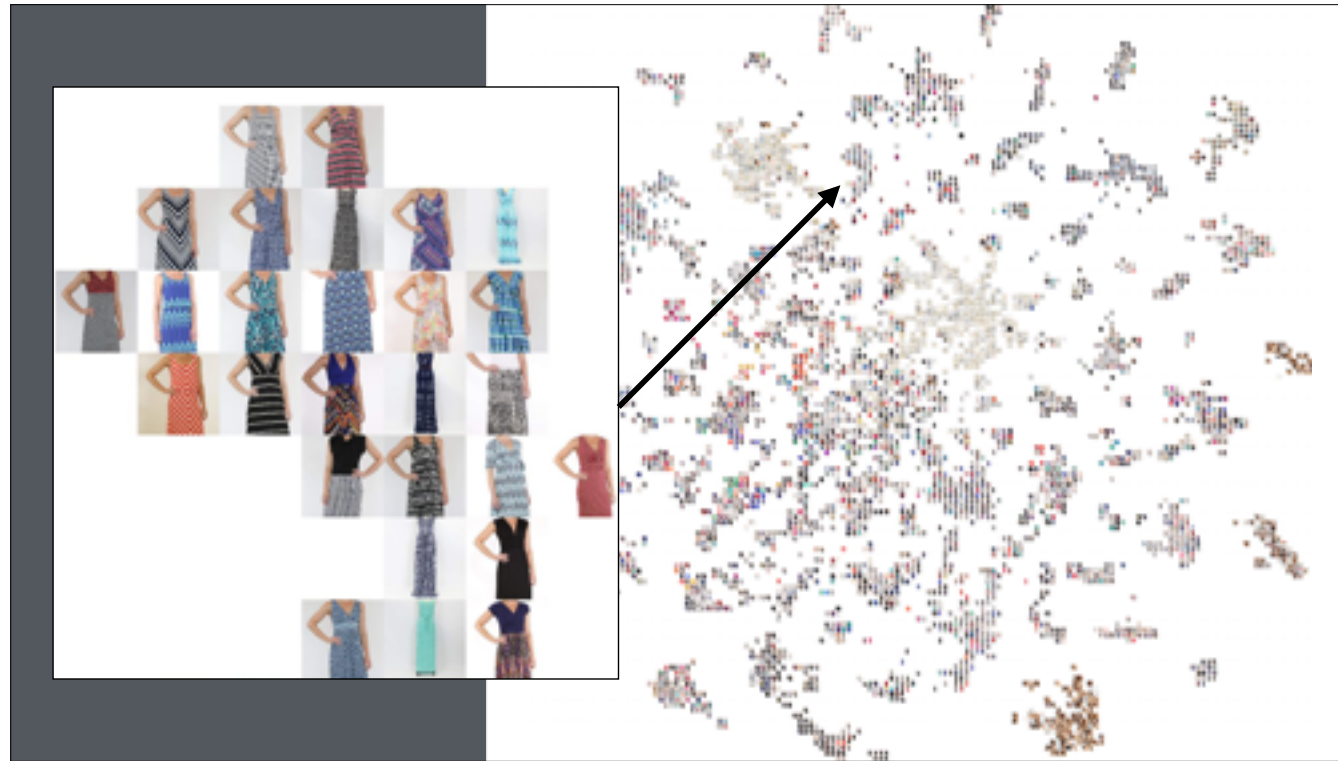


clunky jewelry

dangling delicate jewelry elsewhere



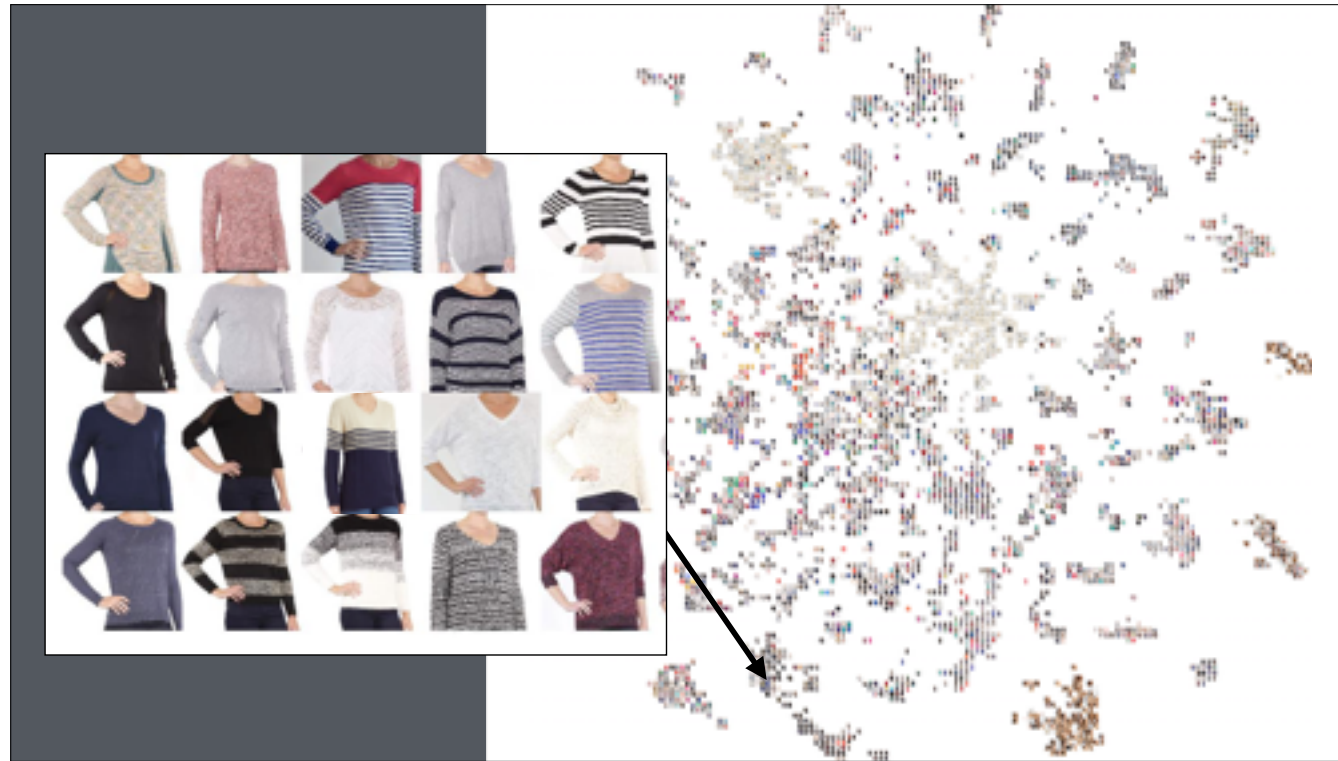
topics on patterns, styles — this cluster is similarly described as high contrast tops with popping colors



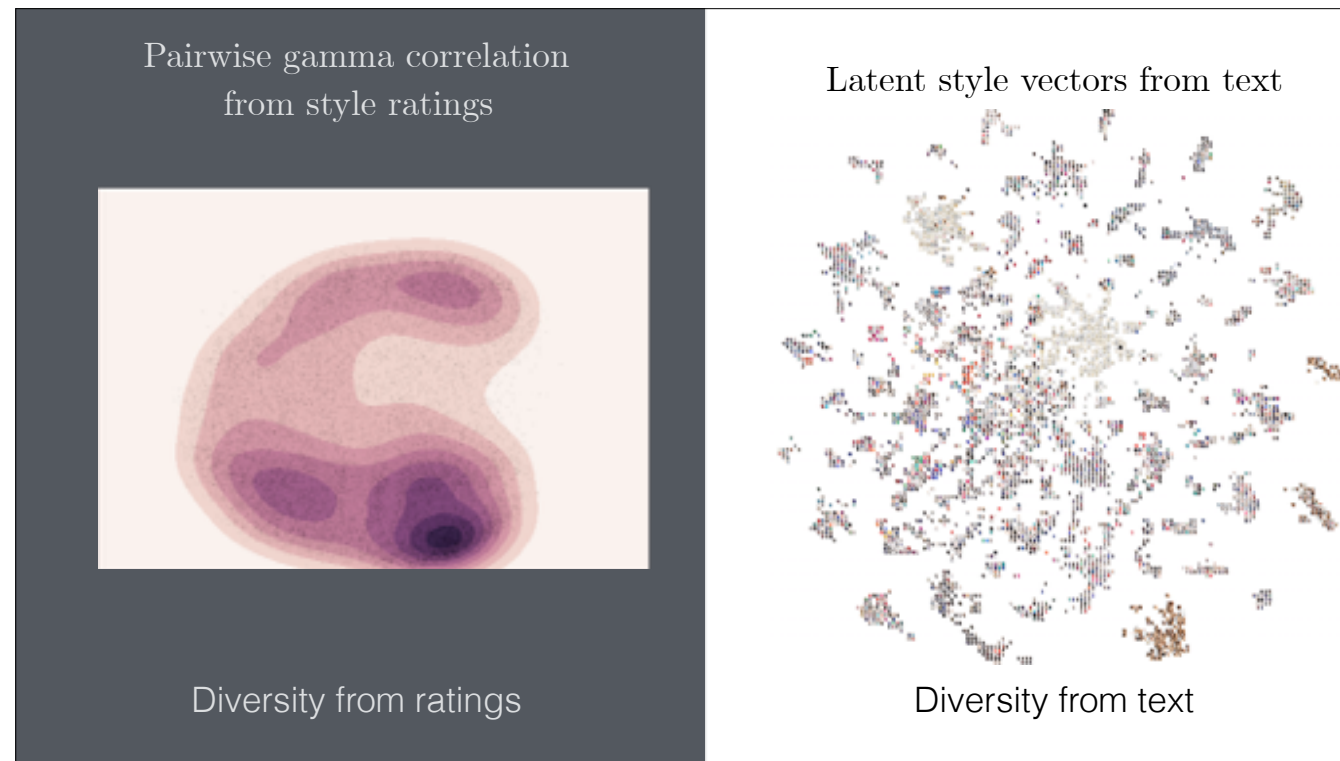
bright dresses for a warm summer



maternity line clothes



not just visual topics, but also topics about fit



Lots of structure in both — but the diversity much higher in the text

Maybe obvious: but the way people describe items is fundamentally richer than the style ratings

lda vs word2vec

“I love finding new designer brands for jeans”



word2vec is *local*:
one word predicts a nearby word

as if the world where one very long text string. no end of documents, no end of sentence, etc.

and a window across words

client_comments

I really like the color of this top and the fit but for suc...
Almost too big. Love the dress though. Going to k...
EVERYTHING about this dress is absolutely PERFE...
This was a Winner to Update my look.... thanks...
Love love love!!! Nothing more to say here.
I love finding new designer brands for jeans. I usual...
Didn't think I'd be too interested in jewelry but t...
Love love love the color, pattern and flowiness!

“I love finding new designer brands for jeans”

But text is usually organized.

The diagram illustrates the process of extracting structured data from unstructured text. On the left, a list of client comments is shown. One comment, 'I love finding new designer brands for jeans. I usual...', is highlighted. An arrow points from this comment to a larger, structured version of the same sentence on the right, which is enclosed in quotation marks. The words 'designer' and 'jeans' are highlighted in green and blue respectively in both versions. Below the structured sentence, the text 'But text is usually organized.' is written.

as if the world where one very long text string. no end of documents, no end of sentence, etc.

client_comments	document_id
I really like the color of this top and the fit but for suc...	5943
Almost too big. Love the dress though. Going to k...	5872
EVERYTHING about this dress is absolutely PERFE...	5951
This was a Winner to Update my look.... thanks...	4017
Love love love!!! Nothing more to say here.	5953
I love finding new designer brands for jeans. I usuall...	7681
Didn't think I'd be too interested in jewelry but t...	3870
Love love love the color, pattern and flowiness!	6286

“I love finding new designer brands for jeans”

But text is usually organized.

as if the world where one very long text string. no end of documents, no end of sentence, etc.

client_comments	document_id
I really like the color of this top and the fit but for suc...	5943
Almost too big. Love the dress though. Going to k...	5872
EVERYTHING about this dress is absolutely PERFE...	5951
This was a Winner to Update my look.... thanks...	4017
Love love love!!! Nothing more to say here.	5953
I love finding new designer brands for jeans. I usual...	7681
Didn't think I'd be too interested in jewelry but I...	3870
Love love love the color, pattern and flowiness!	6286

“I love finding new designer brands for jeans”

doc 7681

In LDA, documents *globally* predict words.

these are client comment which are short, only predict dozens of words

but could be legal documents, or medical documents, 10k words — here the difference between global and local algorithms is much more important

typical word2vec vector

[-0.75, -1.25, -0.55, -0.12, +2.2]

typical LDA document vector

[0%, 9%, 78%, 11%]

typical word2vec vector

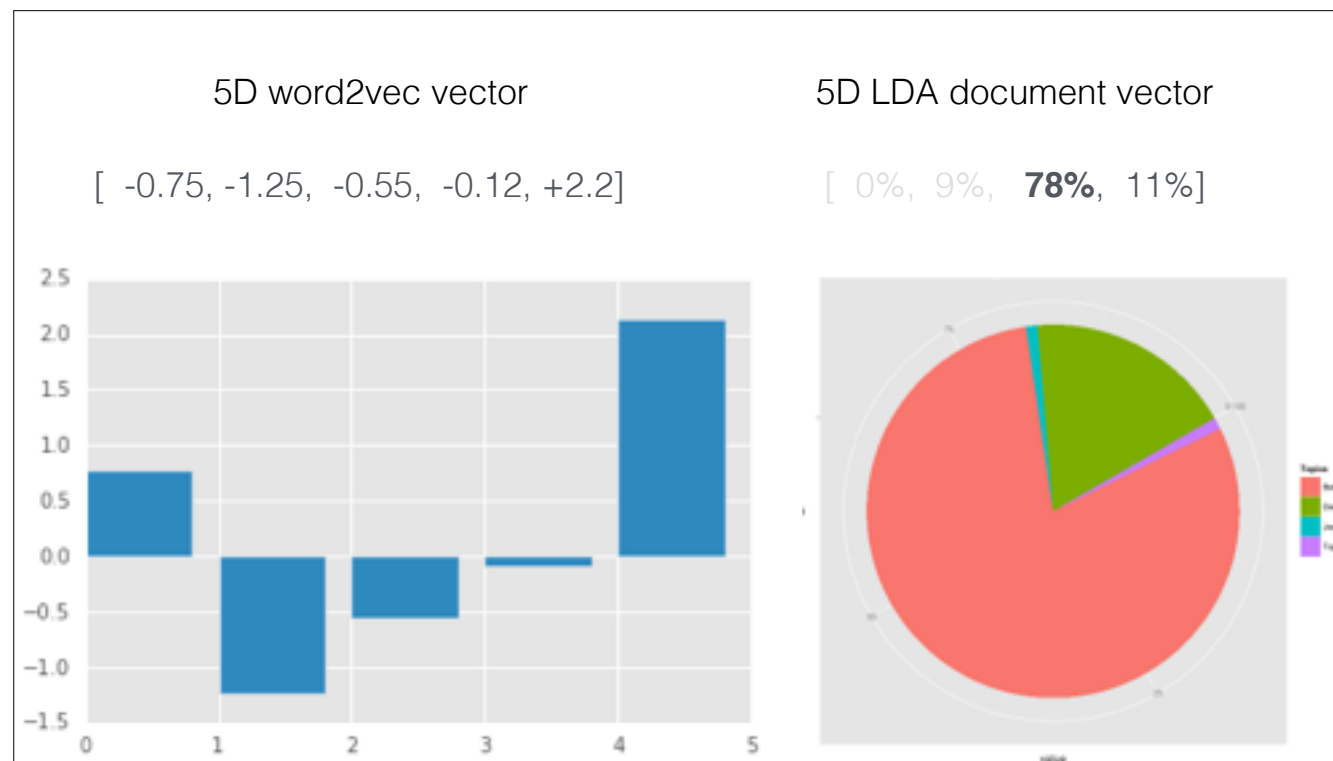
[-0.75, -1.25, -0.55, -0.12, +2.2]

All real values

typical LDA document vector

[0%, 9%, **78%**, 11%]

All sum to 100%



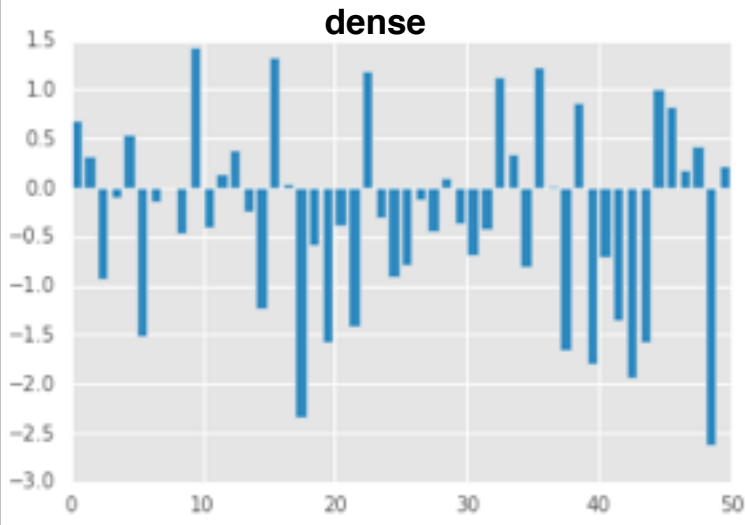
LDA is a *mixture*

w2v is a bunch of real numbers — more like an *address*

much easier to say to another human 78% of something rather than it is +2.2 of something and -1.25 of something else

100D word2vec vector

[-0.75, -1.25, -0.55, -0.27, -0.94, 0.44, 0.05, 0.31 ... -0.12, +2.2]



100D LDA document vector

[0%0%0%0%0% ... 0%, 9%, **78%**, 11%]



100D word2vec vector

[-0.75, -1.25, -0.55, -0.27, -0.94, 0.44, 0.05, 0.31 ... -0.12, +2.2]

Similar in 100D ways
(very **flexible**)

100D LDA document vector

[0%0%0%0%0% ... 0%, 9%, **78%**, 11%]

Similar in fewer ways
(more **interpretable**)

+mixture
+sparse


can we do both? **lda2vec**

series of exp

grain of salt

very new — no good quantitative results only qualitative (but promising!)

The goal:
Use all of this context to learn
interpretable topics.

 [@chrisemoody](#)

client_comments

[Blurred text]

I love finding new designer
brands for jeans. I usua...

Didn't think I'd be too
interested in jewelry but t...

word2vec $\rightarrow P(v_{OUT} | v_{IN})$

Use this at SF.
typical table
w2v will use w-w

The goal:
Use all of this context to learn
interpretable topics.

 @chrisemoody

client_comments	document_id
[REDACTED]	5943
[REDACTED]	5872
[REDACTED]	5951
[REDACTED]	4017
[REDACTED]	5953
I love finding new designer brands for jeans. I usual...	7681
Didn't think I'd be too interested in jewelry but t...	3870
[REDACTED]	6286

word2vec →

LDA →

this document is
80% high fashion

this document is
60% style

$$P(v_{OUT} | v_{DOC})$$

LDA will use that doc ID column
you can use this to steer the business as a whole

The goal:
Use all of this context to learn
interpretable topics.

client_comments	document_id	zip_code
[REDACTED]	5943	52
[REDACTED]	5872	194
[REDACTED]	5951	158
[REDACTED]	4017	991
[REDACTED]	5953	193
I love finding new designer brands for jeans. I usual...	7681	314
Didn't think I'd be too interested in jewelry but t...	3870	43
[REDACTED]	6286	151

word2vec →
LDA →

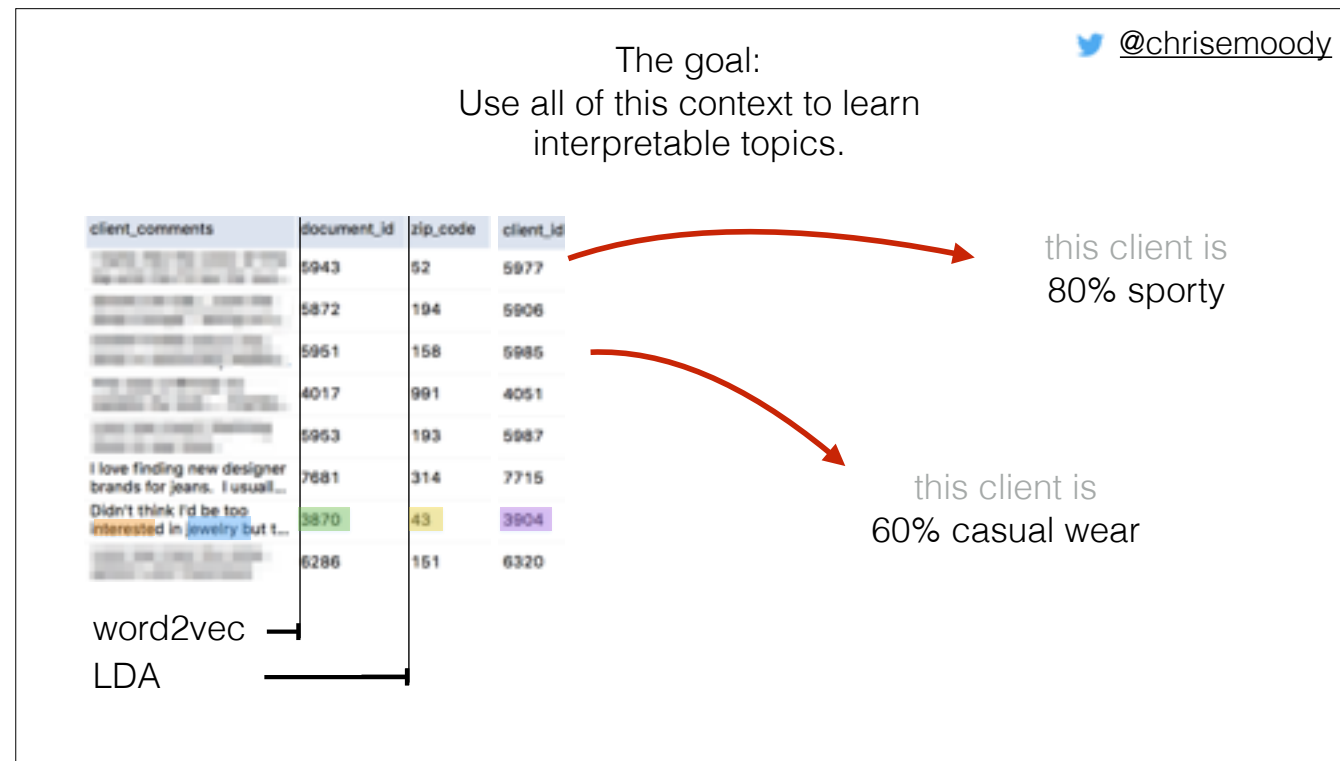
this zip code is
80% hot climate

this zip code is
60% outdoors wear

But doesn't predict word-to-word relationships.

in texas, maybe i want more lonestars & stirrup icons

in austin, maybe i want more bats



love to learn client topics
are there 'types' of clients? q every biz asks

so this is the promise of lda2vec

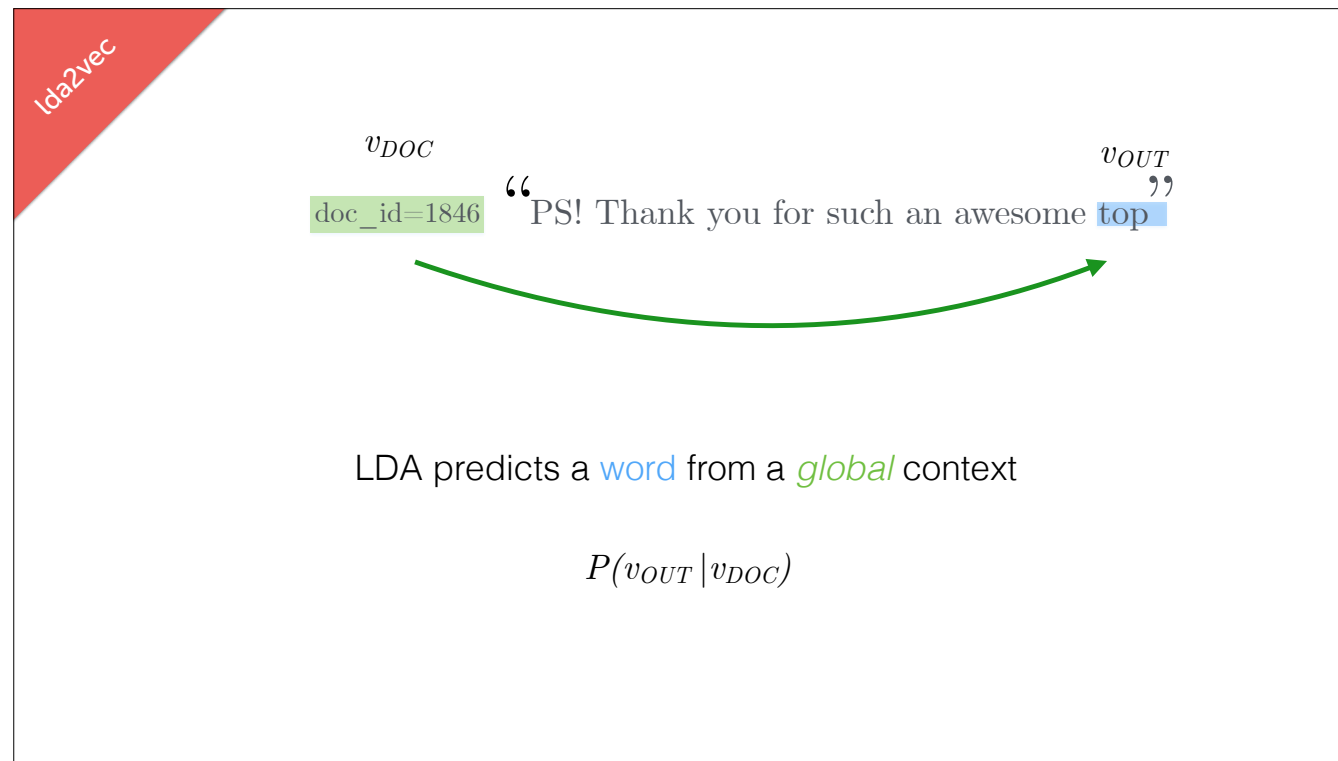
lda2vec

“PS! Thank you for such an v_{IN} awesome v_{OUT} top”

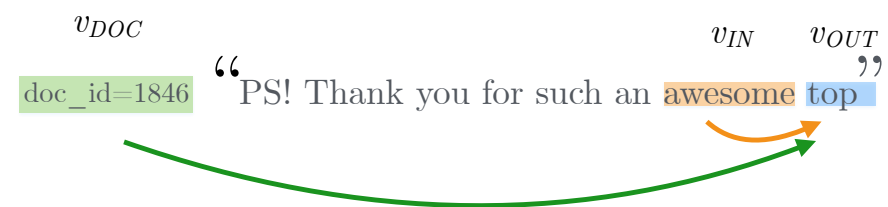
word2vec predicts *locally*:
one word predicts a nearby word

$$P(v_{OUT} | v_{IN})$$

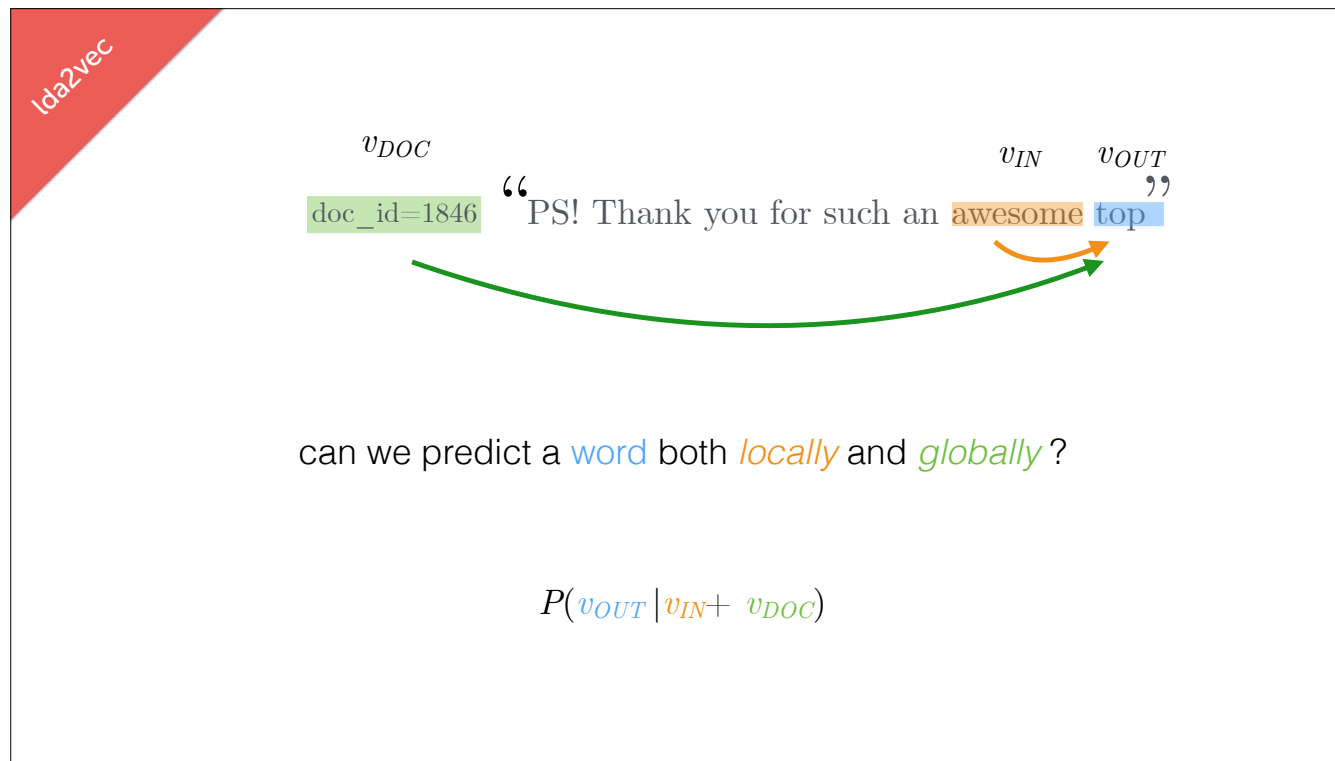
But doesn't predict word-to-word relationships.



But doesn't predict word-to-word relationships.



can we predict a word both *locally* and *globally*?



doc vector captures long-distance dependencies

word vector captures short-distance

v_{DOC} “ PS! Thank you for such an v_{IN} awesome v_{OUT} top ”
doc_id=1846

can we predict a word both *locally* and *globally*?

$$P(v_{OUT} | v_{IN} + v_{DOC})$$

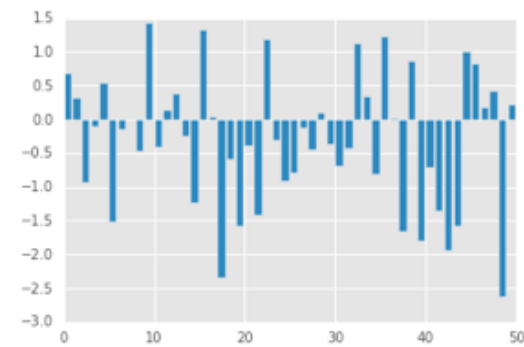
*very similar to the Paragraph Vectors / doc2vec

lda2vec

This works! 😊 But v_{DOC} isn't as interpretable as the LDA topic vectors. 😞

Too many documents. I really like that document X is 70% in topic 0, 30% in topic1, ...

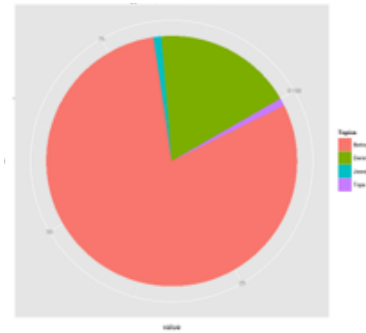
This works! 😊 But v_{DOC} isn't as interpretable as the LDA topic vectors. 😞



Too many documents. I really like that document X is 70% in topic 0, 30% in topic1, ...

about as interpretable a hash

This works! 😊 But v_{DOC} isn't as interpretable as the LDA topic vectors. 😞



Too many documents. I really like that document X is 70% in topic 0, 30% in topic1, ...

This works! 😊 But v_{DOC} isn't as interpretable as the LDA topic vectors. 😞

We're missing *mixtures* & *sparsity*.

Too many documents. I really like that document X is 70% in topic 0, 30% in topic1, ...

This works! 😊 But v_{DOC} isn't as interpretable as the LDA topic vectors. 😞

Let's make v_{DOC} into a mixture...

Too many documents. I really like that document X is 70% in topic 0, 30% in topic1, ...

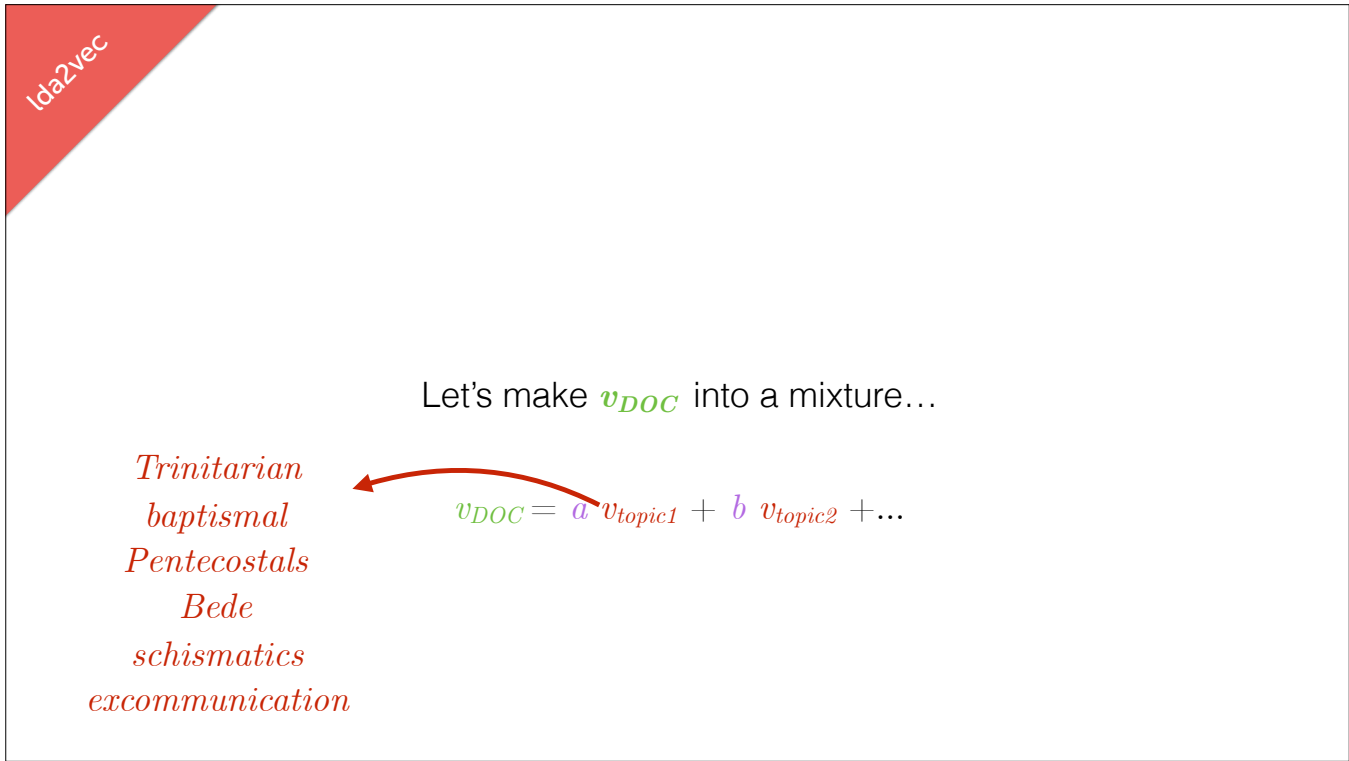
Let's make v_{DOC} into a mixture...

$$v_{DOC} = a \ v_{topic1} + b \ v_{topic2} + \dots \quad (\text{up to } k \text{ topics})$$

sum of other word vectors

intuition here is that 'hanoi = vietnam + capital' and lufthansa = 'germany + airlines'

so we think that document vectors should also be some word vector + some word vector



twenty newsgroup dataset, free, canonical

topic 1 = “religion”

Trinitarian

baptismal

Pentecostals

Bede

schismatics

excommunication

Let's make v_{DOC} into a mixture...

$$v_{DOC} = a v_{topic1} + b v_{topic2} + \dots$$

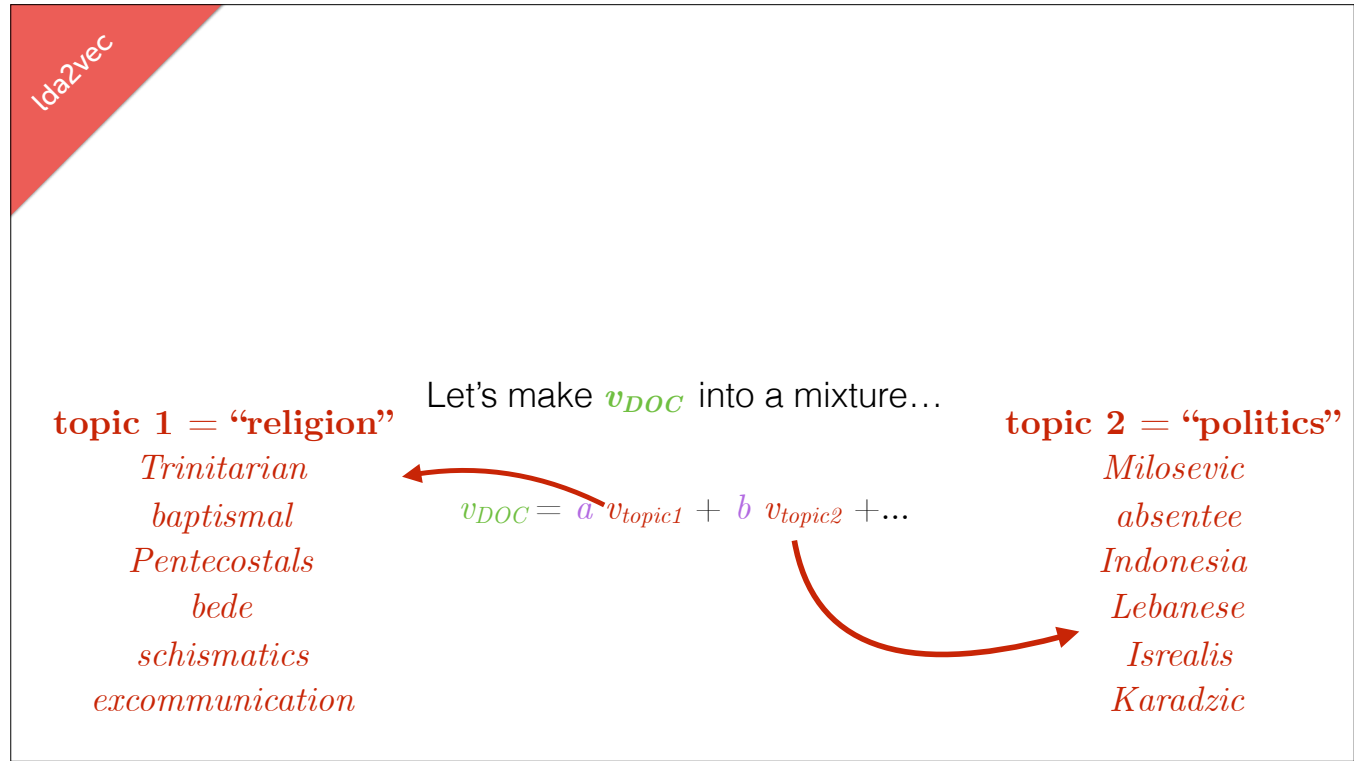
topic 1 = “religion”

Trinitarian
baptismal
Pentecostals
Bede
schismatics
excommunication

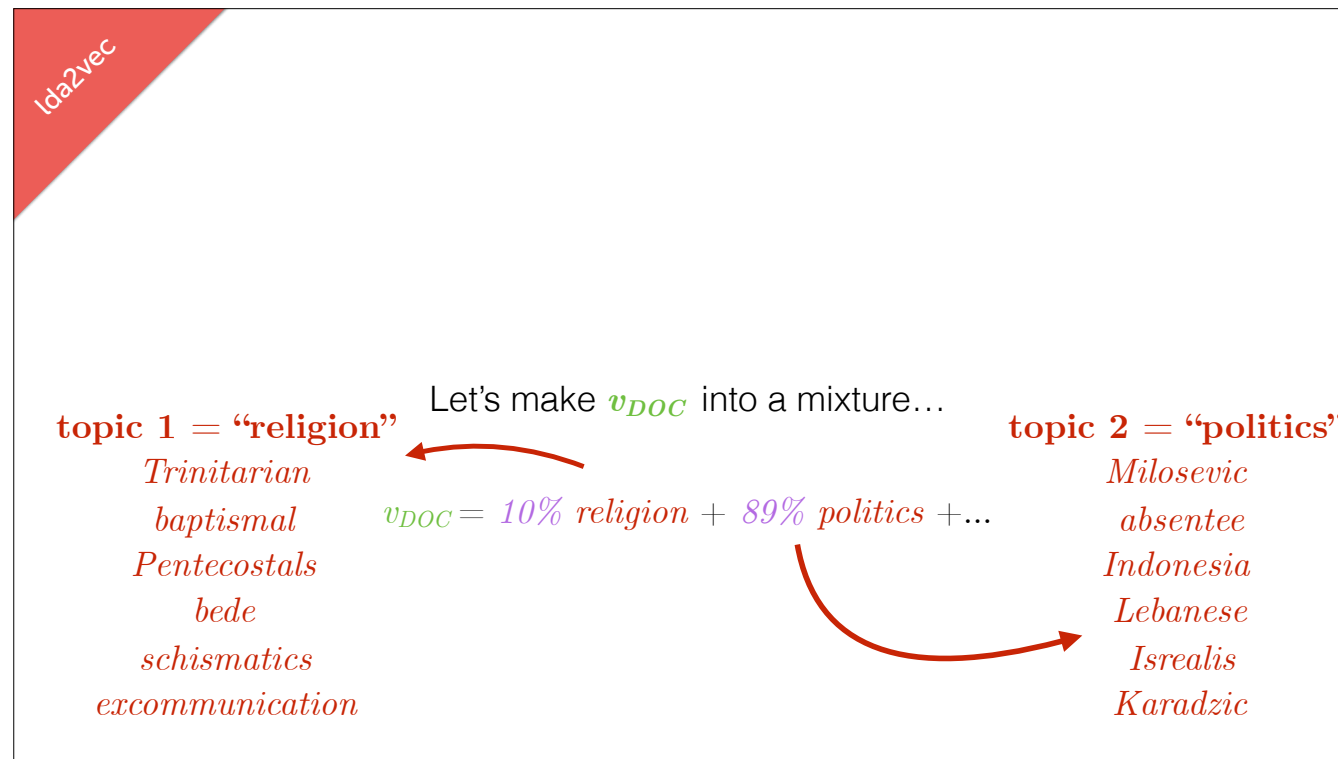
Let's make v_{DOC} into a mixture...

$$v_{DOC} = a v_{topic1} + b v_{topic2} + \dots$$

Milosevic
absentee
Indonesia
Lebanese
Isrealis
Karadzic



purple a,b coefficients tell you how much it is that topic



Doc is now 10% religion 89% politics

mixture models are powerful for interpretability

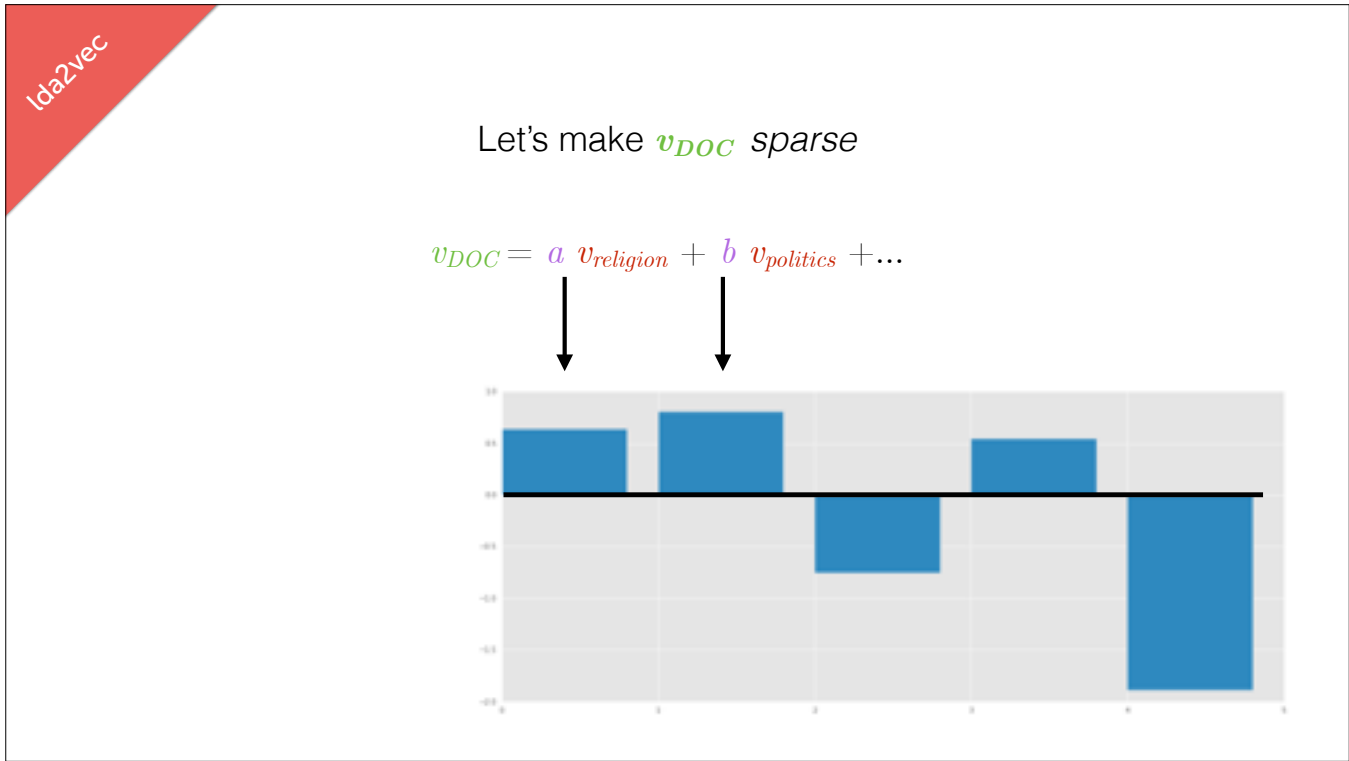
Let's make v_{DOC} *sparse*

$$v_{DOC} = a v_{religion} + b v_{politics} + \dots$$

$$\begin{array}{cc} \downarrow & \downarrow \\ [-0.75, & -1.25, \dots] \end{array}$$

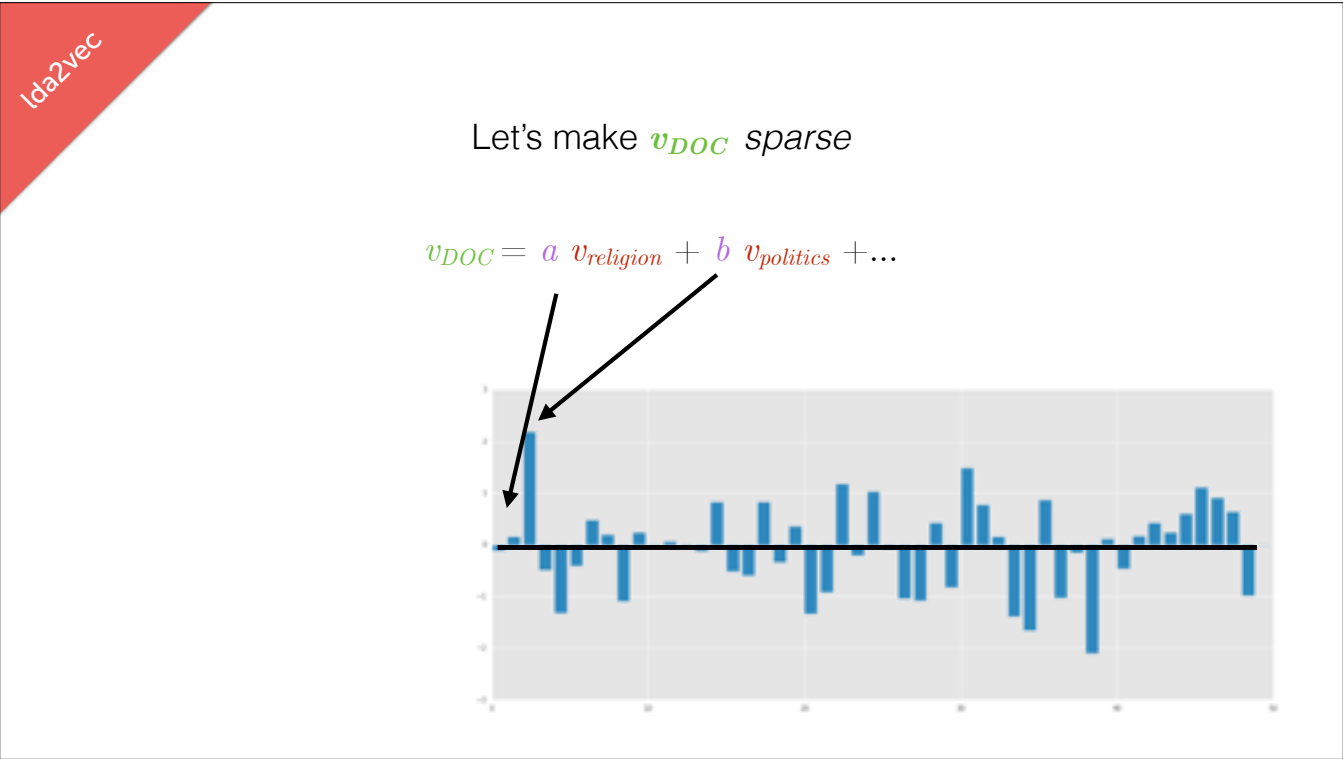
Now 1st time I did this...

Hard to interpret. What does -1.2 politics mean? math works, but not intuitive



How much of this doc is in religion, how much in poltics

but doesn't work when you have more than a few



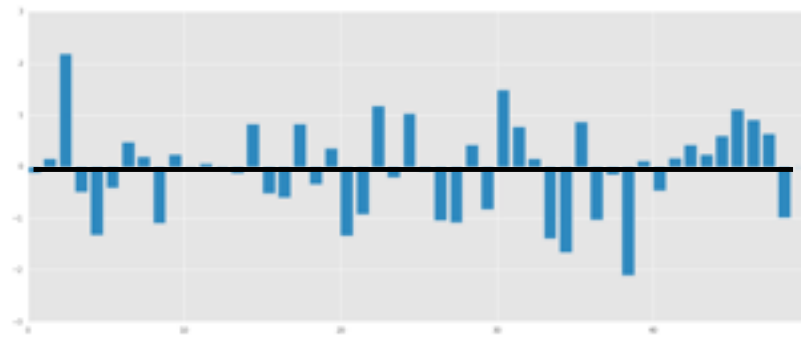
How much of this doc is in religion, how much in cars

but doesn't work when you have more than a few

Let's make v_{DOC} sparse

$$v_{DOC} = a v_{religion} + b v_{politics} + \dots$$

$$\{a, b, c, \dots\} \sim \text{dirichlet}(\alpha)$$



trick we can steal from bayesian

make it dirichlet

skipping technical details

make everything sum to 100%

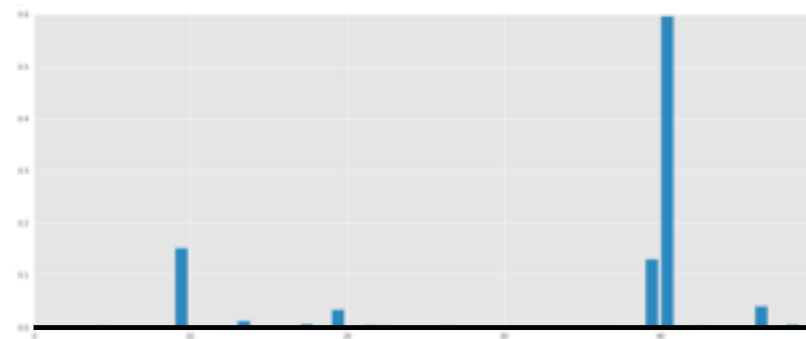
penalize non-zero

force model to only make it non-zero w/ lots of evidence

Let's make v_{DOC} sparse

$$v_{DOC} = a v_{religion} + b v_{politics} + \dots$$

$$\{a, b, c, \dots\} \sim \text{dirichlet}(\alpha)$$



sparsity-inducing effect.

similar to the lasso or l1 reg, but dirichlet

few dimensions, sum to 100%

I can say to the CEO, set of docs could have been in 100 topics, but we picked only the best topics

The goal:
Use all of this context to learn
interpretable topics.

 @chrisemoody

client_comments	document_id
[REDACTED]	5943
[REDACTED]	5872
[REDACTED]	5951
[REDACTED]	4017
[REDACTED]	5953
I love finding new designer brands for jeans. I usual...	7681
Didn't think I'd be too interested in jewelry but t...	3870
[REDACTED]	6286

word2vec

LDA

lda2vec


this document is
80% high fashion

this document is
60% style

$$P(v_{OUT} | v_{IN} + v_{DOC})$$

go back to our problem lda2vec is going to use all the info here

The goal:
Use all of this context to learn
interpretable topics.

 @chrisemoody

client_comments	document_id	zip_code
[REDACTED]	5943	52
[REDACTED]	5872	194
[REDACTED]	5951	158
[REDACTED]	4017	991
[REDACTED]	5953	193
I love finding new designer brands for jeans. I usua...	7681	314
Didn't think I'd be too interested in jewelry but t...	3870	43
[REDACTED]	6286	151

word2vec

LDA

lda2vec

$$P(v_{OUT} | v_{IN} + v_{DOC} + v_{ZIP})$$

add column = adding a term
add features in an ML model

The goal:
Use all of this context to learn
interpretable topics.

client_comments	document_id	zip_code
[REDACTED]	5943	52
[REDACTED]	5872	194
[REDACTED]	5951	158
[REDACTED]	4017	991
[REDACTED]	5953	193
I love finding new designer brands for jeans. I usua...	7681	314
Didn't think I'd be too interested in jewelry but t...	3870	43
[REDACTED]	6286	151

word2vec

LDA

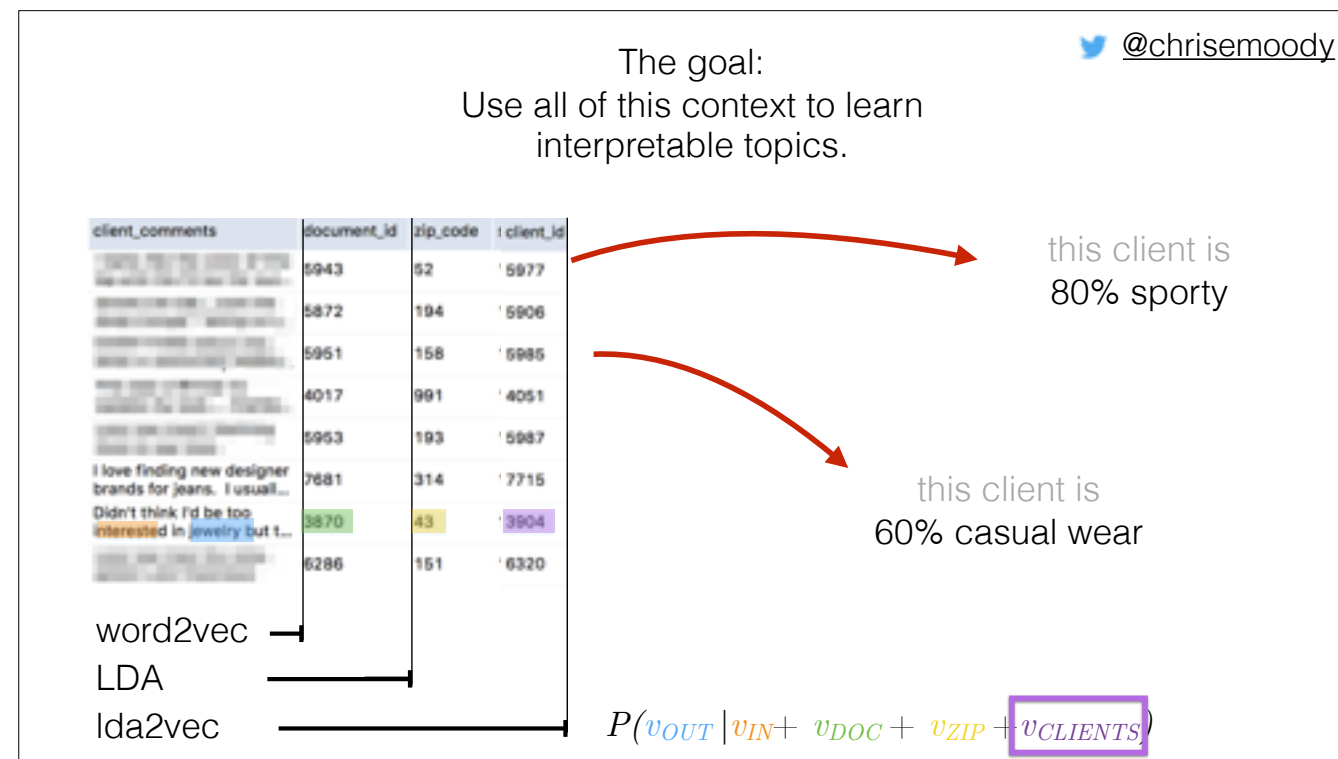
lda2vec

$$P(v_{OUT} | v_{IN} + v_{DOC} + v_{ZIP})$$

this zip code is
80% hot climate

this zip code is
60% outdoors wear

in addition to doc topics, like 'rec SF'



client topics — sporty, casual,
this is where if she says '3rd trimester' — identify a future mother
'scrubs' — medicine

The goal:
Use all of this context to learn
interpretable topics.

 @chrisemoody

client_comments	document_id	zip_code	client_id	sold
...	5943	52	5977	1
...	5872	194	5906	1
...	5951	158	5985	1
...	4017	991	4051	1
...	5953	193	5987	1
I love finding new designer brands for jeans. I usually...	7581	314	7215	1
Didn't think I'd be too interested in jewelry but t...	3870	43	3904	1
...	6286	151	6320	1

Can also make the topics
supervised so that they predict
an outcome.

word2vec

LDA

lda2vec

$$P(v_{OUT} | v_{IN} + v_{DOC} + v_{ZIP} + v_{CLIENTS})$$

$$P(sold | v_{CLIENTS})$$

helps fine-tune topics so that correlate with your favorite business metric

align topics w/ expectations

helps us guess when revenue goes up what the leading causes are

Figure 1 displays the relationship between topic 10 and other topics, showing the top 30 most relevant terms for topic 10 (4.2% of tokens).

The top-left panel shows the intertopic distance map (via multidimensional scaling) for topic 10. The top-right panel shows the top 30 most relevant terms for topic 10 (4.2% of tokens).

The bottom-left panel shows the weighted network graph, illustrating the relationship between topic 10 and other topics. The bottom-right panel shows the overall term frequency, with the top 30 most relevant terms for topic 10 highlighted in red.

Top 30 Most Relevant Terms for Topic 10 (4.2% of tokens)

Term	Frequency (approx.)
introduction	10
positive	10
discovery	12
relationships	10
general	10
regulation	10
production	12
genes	10
function	12
cell	12
biology	15
study	15
results	15
environment	18
data	10
model	15
clinical	10
pathways	10
transcript	20
network	25
particular	10
development	10
proteins	10
approach	10
framework	10
pathway	10
top	25

Overall term frequency

Red bars represent the frequency of terms within the selected topic. Blue bars represent the overall frequency of terms.

Legend:

- Red bar: Frequency within the selected topic
- Blue bar: Overall frequency

Figure 1 displays the relationship between topic 10 and other topics, showing the top 30 most relevant terms for topic 10 (4.2% of tokens).

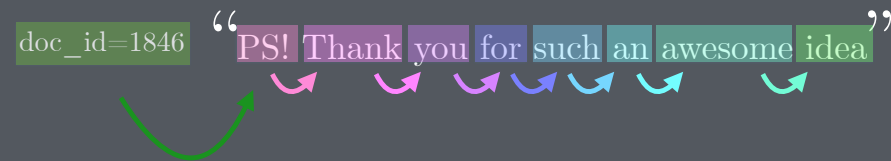
github.com/cemoody/lda2vec



@chrisemoody

Can we model topics to sentences?

lda2lstm

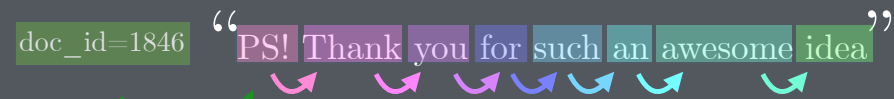


SF is all about mixing cutting edge algorithms but we absolutely need interpretability. human component to algos is not negotiable

Could we demand the model make us a sentence that is 80% religion, 10% politics?

classify word level, LSTM on sentence, LDA on document level

doc_id=1846 “PS! Thank you for such an awesome idea”

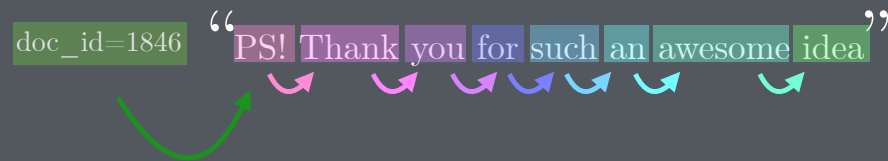


Can we represent the internal LSTM
states as a dirichlet mixture?

Dirichlet-squeeze internal states and manipulations, that maybe will help us understand the science of LSTM dynamics — because seriously WTF is going on there

Can we model topics to sentences?

lda2lstm



Can we model topics to images?

lda2ae



TJ Torres



Can we also extend this to image generation? TJ is working on a ridiculous VAE/GAN model... can we throw in a topic model? Can we say make me an image that is 80% sweater, and 10% zippers, and 10% elbow patches?



?

 @chrisemoody

Multithreaded
Stitch Fix






Bonus slides







Paragraph Vectors
(Just extend the context window)

Content dependency
(Change the window grammatically)

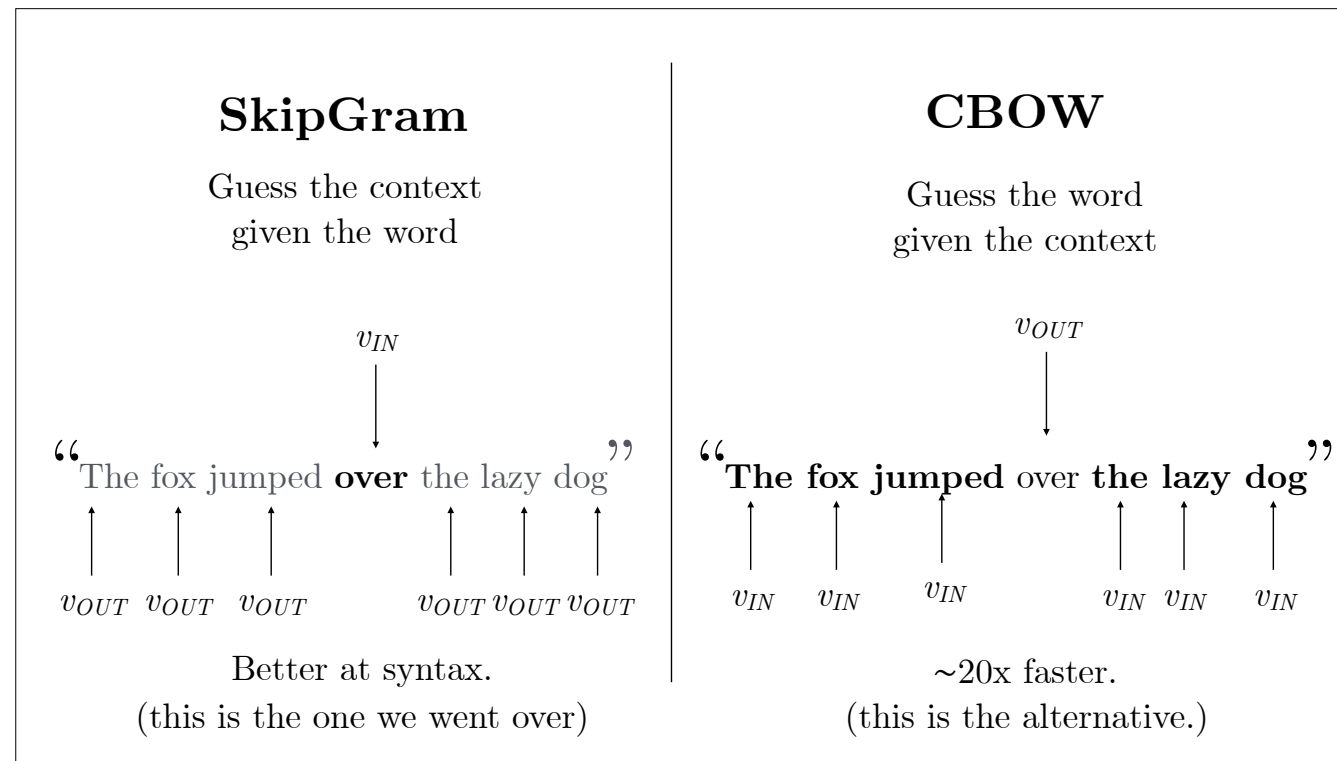
Social word2vec (deepwalk)
(Sentence is a walk on the graph)

Spotify
(Sentence is a playlist of song_ids)

Stitch Fix
(Sentence is a shipment of five items)

Relationship	Example 1	Example 2	Example 3
France - Paris	Italy: Rome	Japan: Tokyo	Florida: Tallahassee
big - bigger	small: larger	cold: colder	quick: quicker
Miami - Florida	Baltimore: Maryland	Dallas: Texas	Kona: Hawaii
Einstein - scientist	Messi: midfielder	Mozart: violinist	Picasso: painter
Sarkozy - France	Berlusconi: Italy	Merkel: Germany	Koizumi: Japan
copper - Cu	zinc: Zn	gold: Au	uranium: plutonium
Berlusconi - Silvio	Sarkozy: Nicolas	Putin: Medvedev	Obama: Barack
Microsoft - Windows	Google: Android	IBM: Linux	Apple: iPhone
Microsoft - Ballmer	Google: Yahoo	IBM: McNealy	Apple: Jobs
Japan - sushi	Germany: bratwurst	France: tapas	USA: pizza

See previous



CBOW sums words vectors, loses the order in the sentence

Both are good at semantic relationships

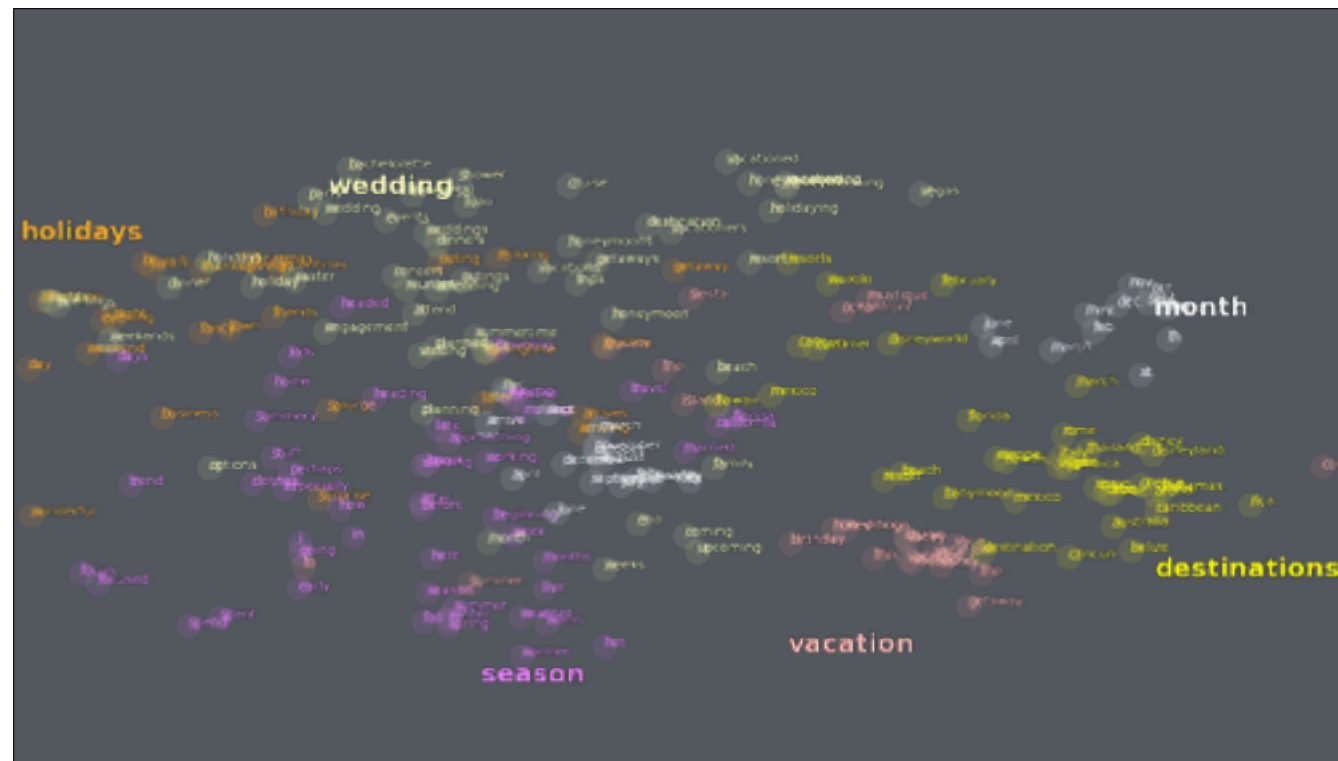
Child and kid are nearby

Or gender in man, woman

If you blur words over the scale of context — 5ish words, you lose a lot grammatical nuance

But skipgram preserves order

Preserves the relationship in pluralizing, for example



Shows that are many words similar to vacation actually come in lots of flavors

- wedding words (bachelorette, rehearsals)
- holiday/event words (birthdays, brunch, christmas, thanksgiving)
- seasonal words (spring, summer,)
- trip words (getaway)
- destinations

LDA Results

Great Stylist

Perfect

I loved every choice in this fix!! Great job!

History

LDA Results

Body Fit

My measurements are 36-28-32. If that helps.
I like wearing some clothing that is fitted.
Very hard for me to find pants that fit right.

History

LDA Results

Sizing

Excited for next

Really enjoyed the experience and the
pieces, sizing for tops was too big.
Looking forward to my next box!

History

LDA Results

Almost Bought

Perfect

It was a great fix. Loved the two items I kept and the three I sent back were close!

History

What I didn't mention

A lot of text (only if you have a specialized vocabulary)

Cleaning the text

Memory & performance

Traditional databases aren't well-suited

False positives

hundreds of millions of words, 1,000 books, 500,000 comments, or 4,000,000 tweets

high-memory and high-performance multicore machine.

Training can take several hours to several days but shouldn't need frequent retraining.

If you use pretrained vectors, then this isn't an issue.

Databases. Modern SQL systems aren't well-suited to performing the vector addition, subtraction and multiplication searching in vector space requires. There are a few libraries that will help you quickly find the most similar items¹²: annoy, ball trees, locality-sensitive hashing (LSH) or FLANN.

False-positives & exactness. Despite the impressive results that come with word vectorization, no NLP technique is perfect. Take care that your system is robust to results that a computer deems relevant but an expert human wouldn't.

and now for something **completely crazy**

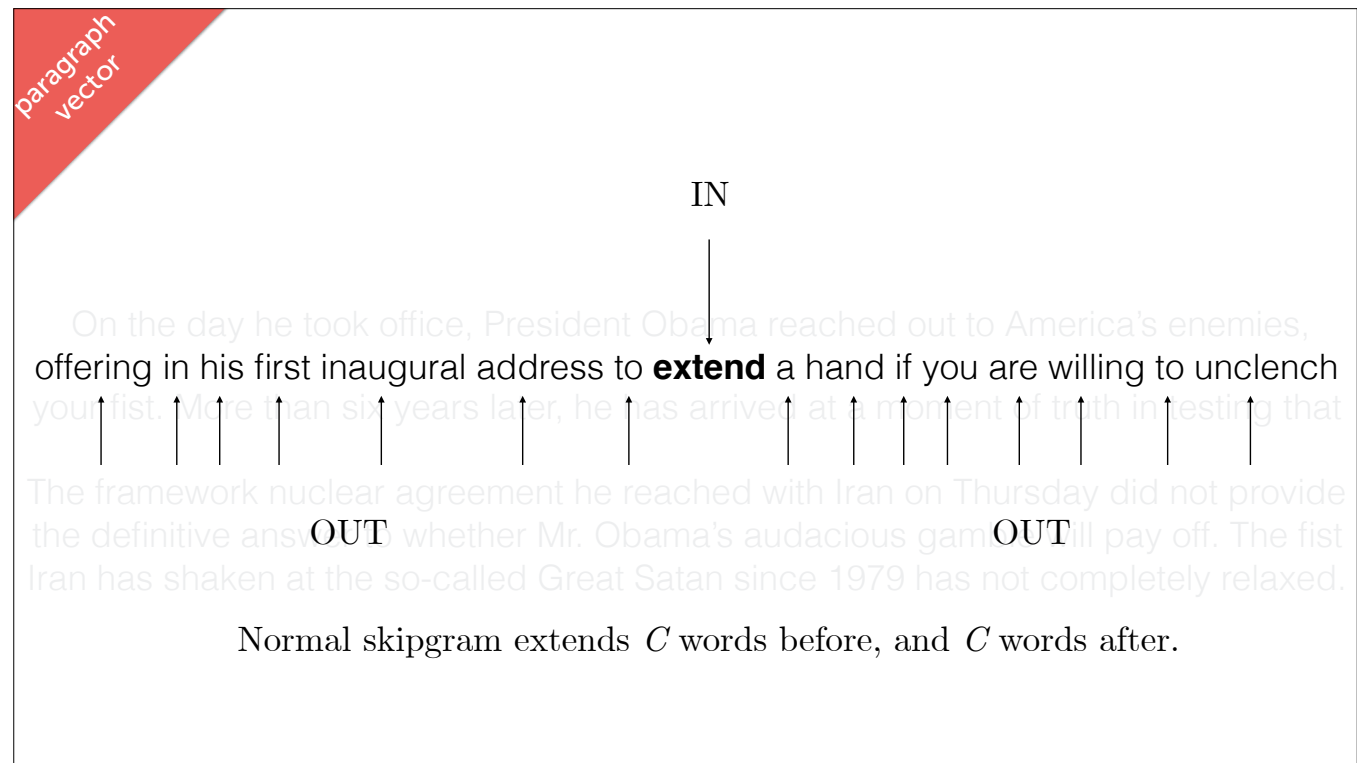
All of the following ideas will change what
'words' and 'context' represent.

But we'll still use the same w2v algo

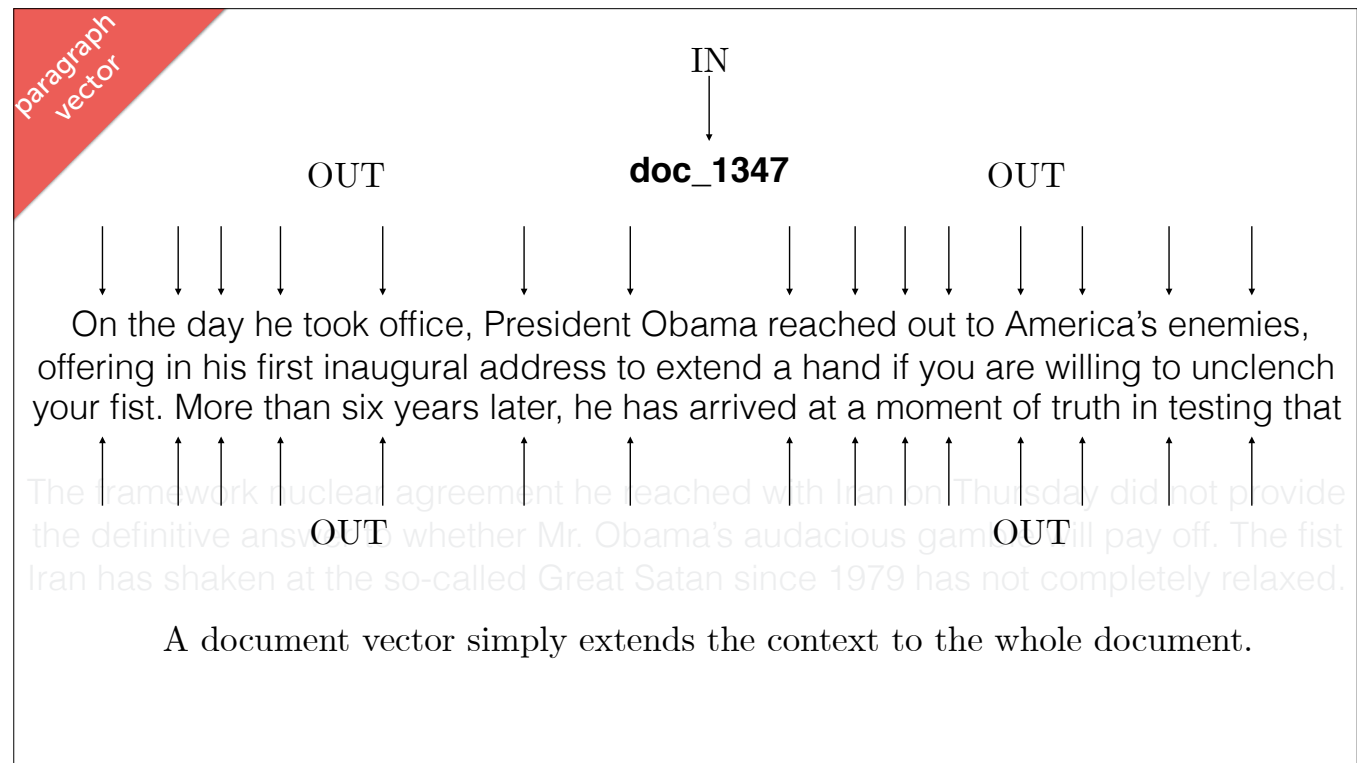
paragraph
vector

What about summarizing documents?

On the day he took office, President Obama reached out to America's enemies, offering in his first inaugural address to **extend** a hand if you are willing to unclench your fist. More than six years later, he has arrived at a moment of truth in testing that



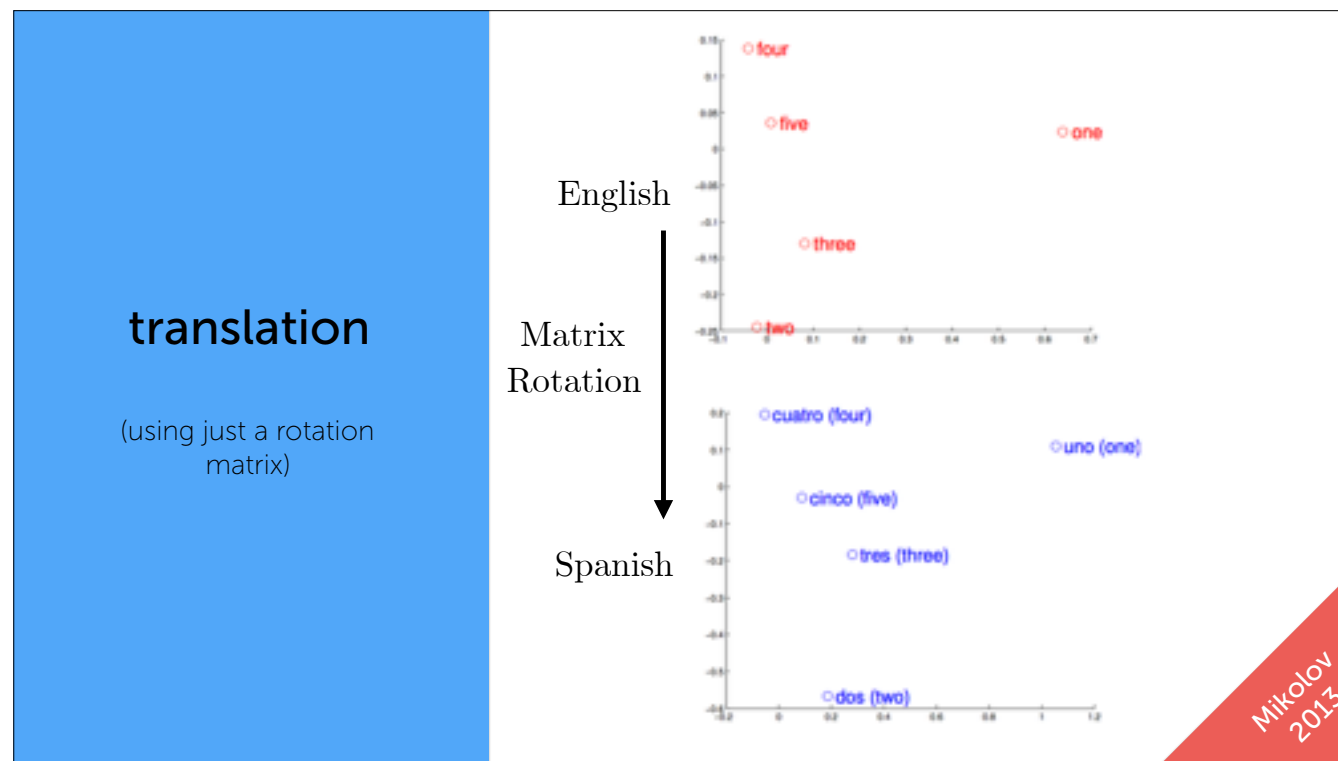
Except we stay inside a sentence



```
from gensim.models import Doc2Vec
fn = "item_document_vectors"
model = Doc2Vec.load(fn)
model.most_similar('pregnant')
matches = list(filter(lambda x: 'SENT_' in x[0], matches))

# ['...I am currently 23 weeks pregnant...',
#  '...I'm now 10 weeks pregnant...',
#  '...not showing too much yet...',
#  '...15 weeks now. Baby bump...',
#  '...6 weeks post partum!...',
#  '...12 weeks postpartum and am nursing...',
#  '...I have my baby shower that...',
#  '...am still breastfeeding...',
#  '...I would love an outfit for a baby shower...']
```

sentence
search



Blows my mind

Explain plot

Not a complicated NN here

Still have to learn the rotation matrix — but it generalizes very nicely.

Have analogies for every linalg op as a linguistic operator: + and - and matrix multiplies

Robust framework and new tools to do science on words

context
dependent

Australian scientist **discovers** star with telescope

context \pm 2 words

Livy & Goldberg
2014

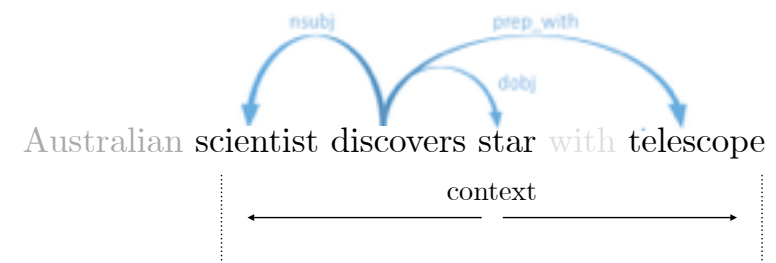
context
dependent



Livy & Goldberg
2014

What if we

context
dependent



Ley & Goldberg
2014

context
dependent

	BoW	DEPS
hogwarts	dumbledore hallows half-blood malfoy snape	sunnydale collinwood calarts greendale millfield
topically-similar	vs	'functionally' similar

Levy & Goldberg
2014

context
dependent

Also show that SGNS is simply factorizing:

$$w * c = PMI(w, c) - \log k$$

This is **completely** amazing!

Intuition: positive associations (canada, snow)
stronger in humans than negative associations
(what is the opposite of Canada?)

Ley & Goldberg
2014

Also means we can do SVD-like techniques to get a convex $w2v$, uses fast lining libs, uses compressed word count matrix so also better storage.... but not online

word2vec

learn word vectors from
sentences

“The fox jumped over **the** lazy dog”

v_{OUT} v_{OUT} v_{OUT} v_{OUT} v_{OUT} v_{OUT}

deepwalk

‘words’ are graph vertices
‘sentences’ are random walks on the
graph

$v_{46} \rightarrow v_{45} \rightarrow v_{71} \rightarrow v_{24} \rightarrow v_5$

Perozzi
et al 2014

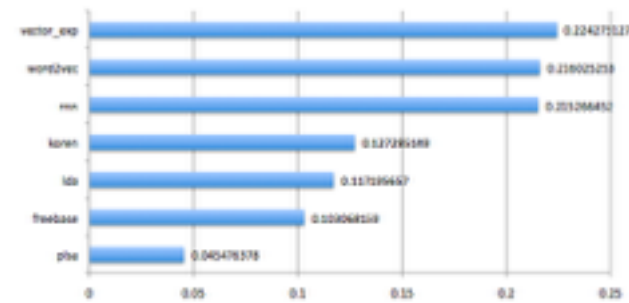
Playlists at Spotify

‘words’ are songs
‘sentences’ are playlists

sequence
learning

Playlists at Spotify

Great performance on ‘related artists’



Erik
Bernhardsson

Fixes at Stitch Fix

Let's try:

'words' are styles

'sentences' are fixes

sequence
learning

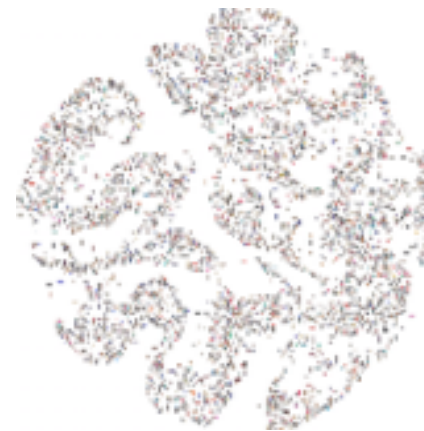
Fixes at Stitch Fix

Learn similarity between styles
because they co-occur

Learn 'coherent' styles

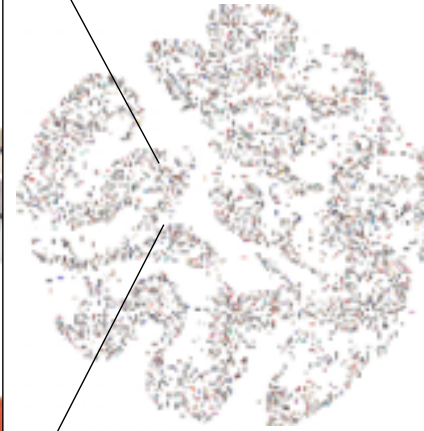
sequence
learning

Fixes at
Stitch Fix?



Got lots of structure!

sequence
learning



sequence
learning







A specific lda2vec model

Our text blob is a comment that comes from a region_id and a style_id





$$L = \sigma(c * w) + \sigma(-c * w_{neg})$$

$$context = c_{ij} = \vec{region}_i + style_j$$

$$region_i = \sum_{k=0}^{n_{topics}} u_{ik} \cdot \vec{m}_k$$

$$style_j = \sum_{l=0}^{n_{topics}} u_{jl} \cdot \vec{n}_l$$

$$\vec{u} \sim dirichlet(\alpha_1)$$

$$\vec{v} \sim dirichlet(\alpha_2)$$

$$take_rate_in_region \sim 5.0 * \sigma(W \cdot \vec{u})$$

The full likelihood model

$$L = \sigma(c * w) + \sigma(-c * w_{neg})$$

$$context = c_{ij} = region_i + style_j$$

$$region_i = \sum_{k=0}^{n_topics} u_{ik} \cdot \vec{m}_k$$

$$style_j = \sum_{l=0}^{n_topics} u_{jl} \cdot \vec{n}_l$$

$$\vec{u} \sim \text{dirichlet}(\alpha_1)$$

$$\vec{v} \sim \text{dirichlet}(\alpha_2)$$

$$take_rate_in_region \sim 5.0 * \sigma(W \cdot \vec{u})$$

$$L = \sigma(\mathbf{c} + \mathbf{w}) + \sigma(-\mathbf{c} + \mathbf{w}_{\text{neg}})$$

First part of the loss function is given **context** predict **word**.

Don't predict a **negative word**. These are words that are in our vocabulary somewhere, but not in our example.

We get negative samples **not** uniformly, but proportional to the word frequency^{3/4} (yes, the 3/4 power is weird and ad hoc but totally works awesomely for word2vec)

$$L = \sigma(c * w) + \sigma(-c * w_{neg})$$

$$context = c_d = region_d + style_d$$

Context is made up from more than one part -- many 'contexts' available.

In this case, instead of one document, we can have many regions, or styles.

In LDA, this context is a single term: the latent document vector that 'generates' words.

In word2vec, this context is the 'pivot' word. Word2vec picks a random 'context' word in the corpus, centers a window around it, and tries to predict other words within that context.

In both word2vec and LDA context is one term, either a document or a word. For lda2vec, we can more than one term, we can have as many contexts as we like!

$$L = \sigma(v \cdot w) + \sigma(-v \cdot w_{neg})$$

$$context = \vec{c}_j = region_j + style_j$$

$$region_j = \sum_{k=0}^{n_topics} u_{jk} \cdot \vec{m}_k$$

$$style_j = \sum_{l=0}^{n_topics} u_{jl} \cdot \vec{n}_l$$

Each context (e.g., **region** or **style**) is decomposed into **topics vectors** and **weights** on those common **topics vectors**. One context has one shared set of topic vectors (think of these as cluster centroids) and every 'document' in that context (think of 1 of 50 states, 1 of 20k styles) has a weight/membership onto each of those topic vectors (think topics like northeast, midwest for region or tops, bottoms, boho, romantic for style topics)

This forces the context vectors onto **a limited set of basis vectors**. Interpret this set, and you can generalize what each region vector and style vector means. For example, one **topics vector** might be close to the **word vector** for 'hand_bag', 'purse', 'bag' indicating that that topic is a handbags topic. And then anything with big **weight** in that topic might be a handbag.

$$L = \sigma(\vec{v} \cdot \vec{w}) + \sigma(-\vec{v} \cdot \vec{w}_{neg})$$

$$context = \vec{c}_j = region_j + style_j$$

$$region_j = \sum_{k=0}^{n_topics} u_{jk} \cdot \vec{m}_k$$

$$style_j = \sum_{l=0}^{n_topics} u_{jl} \cdot \vec{n}_l$$

$$\vec{u} \sim \text{dirichlet}(\alpha_1)$$

$$\vec{v} \sim \text{dirichlet}(\alpha_2)$$

But the weights can still end up being very dense -- which meant everyone of my documents was a mixture of almost every component. This made it difficult to interpret what the document was, because it had membership in many groups.

So next we enforce a simplex with dirichlet & enforce sparsity with the concentration on the **weights**. The dirichlet is also nice but not critical, we could've had a non-negative decomposition or just stuck with all reals. But since Dirichlet components sum to 100%, it is easier to explain to analysts that a document is "10% of some_topic + 90% some_other_topic" rather than saying "-2.3 * some_topic and +0.5 of some_other_topic".

$$L = \sigma(\vec{v} \cdot \vec{w}) + \sigma(-\vec{v} \cdot \vec{w}_{neg})$$

$$context = \vec{c}_j = region_j + style_j$$

$$region_j = \sum_{k=0}^{n_topics} u_{jk} \cdot \vec{m}_k$$

$$style_j = \sum_{l=0}^{n_topics} u_{jl} \cdot \vec{n}_l$$

$$\vec{u} \sim \text{dirichlet}(\alpha_1)$$

$$\vec{v} \sim \text{dirichlet}(\alpha_2)$$

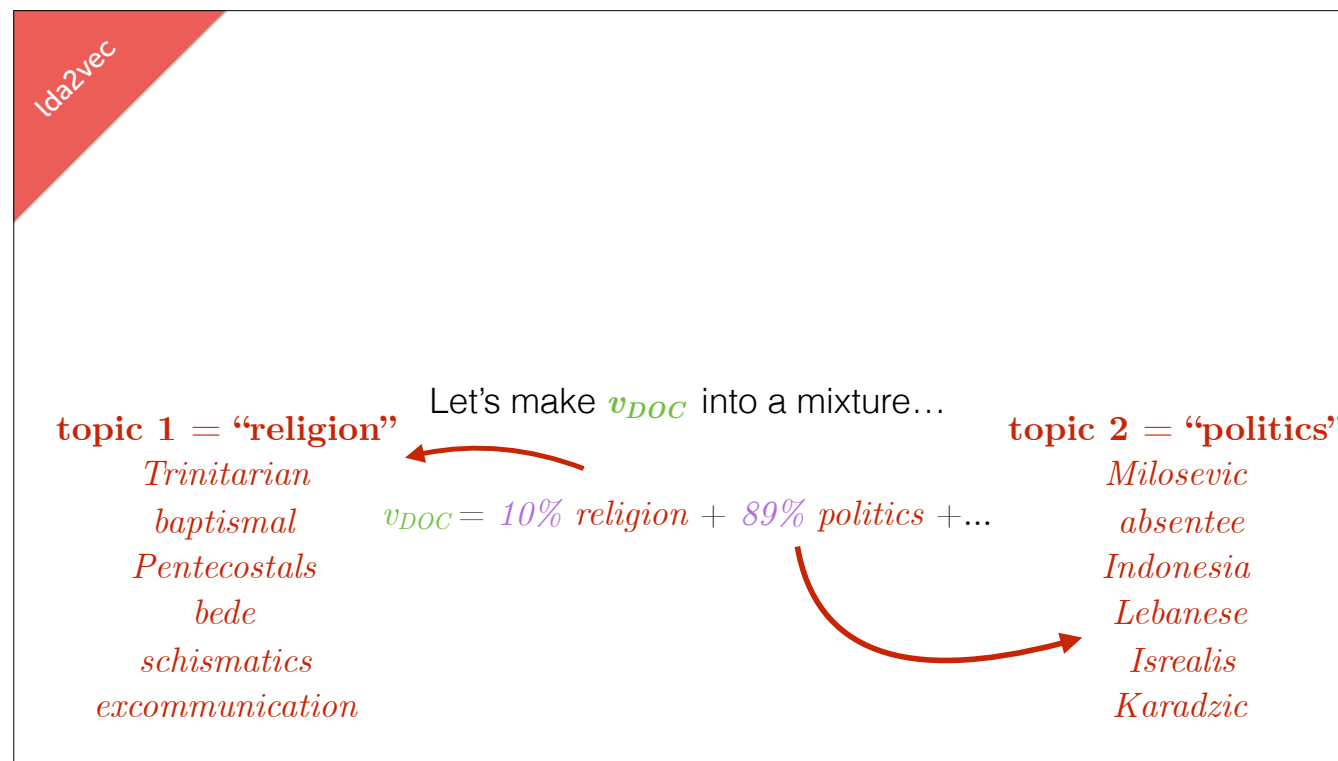
$$take_rate_in_region \sim 5.0 * \sigma(\vec{W} \cdot \vec{u})$$

Finally, we can make this 'supervised' by saying that the topic weights correlate through (matrix \vec{W}) with some target outcome.

Can measure similarity between topic vectors m and n, and word vectors w

This gets you the 'top' words in a topic, can figure out what that topic is





This is now on the 20 newsgroups dataset...

Doc is now 10% religion 89% politics

mixture models are powerful for interpretability